

# Introduction

CS503: Operating systems, Spring 2019

Pedro Fonseca  
Department of Computer Science  
Purdue University


# Hi everyone!

- Campus students and online students!
  - Fancy room, with video & audio recording for online students
- Please do **ask questions**, interactiveness is important


# Finding a photo...today




# Finding a photo...today



# Finding a photo...today




# Finding a photo...today



**Incredibly complex!!!**

# Finding a photo...today



**Incredibly complex!!!**

# Finding a photo...today



Many of these problems have been first found  
in the context of operating systems



# Why take CS503?

- Some of you are going to **write or modify operating systems**
- Some of you will do **research on operating systems**
- Some of you (All?) will **write applications** that interact with operating systems
- Some of you will **leverage operating systems concepts** in other contexts

# Goals for today




- What makes operating systems so exiting?
- Examples of operating system designs
- How does this class operate?
- What is an operating system?

# **A very brief history of operating systems**

## 40's - 50's: no OS

- Machines very expensive
- No high-level PL and no OS
- Programmed machines with patch cables or punched-cards
- IO through punched-cards




**ENIAC Computer  
1945-1955**

[illegible]

***“By 1937... IBM had 32 presses at work in Endicott, N.Y., printing, cutting and stacking 5 to 10 million punched cards every day.”  
— "IBM Archive: Endicott card manufacturing"***

# How many computers do we need?

- *"I think there is a world market for about five computers."* — Thomas J Watson, IBM's president, early 1940s
- Watson's legendary misjudgment did not prove fatal to IBM



## 50's: batch processing

- The run queue is a deck of cards
- Not practical, slow
- One program at a time, no computation during IO
- OS functions: IO handler and loader



# 60's: multiprogrammed batch systems

- When job waits for IO, switch to another job
- New OS functions:
  - Job scheduling policies
  - Memory management
  - Handle asynchronous IO devices
- Computer still not interactive



## 60's-70's: timesharing machines

- Jobs are quickly switched to give appearance of dedicated machine
- New OS functions: preemption, advanced scheduling policies

# 80's: Personal computers

- Entire machine is inexpensive
- One user, one machine
- OS functions were simplified (e.g., DOS):
  - Single user, no timesharing, no virtual memory

# Other recent developments....

- Internet:
  - Network stack became part of every OS
  - Very efficient to support high-throughput and low latencies
- Computers are ubiquitous and come in all forms



# Other recent developments....

## Internet of Things (IoT)




- Constrained devices  
(e.g., limited memory, hw protection)
- May not use a “normal” desktop OSs
- Reliability and security are critical




# How many computers do we need?

- *"I think there is a world market for about five computers."* — Thomas J Watson, IBM's president, early 1940s
- Watson's legendary misjudgment did not prove fatal to IBM



# How many ~~computers~~ do we need? datacenters

- *"I think there is a world market for about five computers."* — Thomas J Watson, IBM's president, early 1940s
  - Watson's legendary misjudgment did not prove fatal to IBM



# Progression over the decades

- Hardware:
  - Vacuum Tubes (1950s) - one bit on the size of a thumb;
  - Transistors (1950s and 1960s) - one bit on the size of a fingernail;
  - Integrated Circuits (1960s and 70s) - thousands of bits on the size of a hand
  - Silicon computer chips (1970s and on) - millions of bits on the size of a finger nail.
- Ease of use and program:
  - Almost impossible to use except by very patient geniuses (1950s);
  - Programmable by highly trained people only (1960s and 1970s);
  - Useable by just about anyone (1980s and on).
- Complexity
- Reliability

# Multics OS, an ambitious OS

- Multics OS (MIT + GE, 1969-):
  - Some say very successful, others very unsuccessful
  - Influential (Linux, Windows)
- Notable properties and functions:
  - Files and memory are treated identically
  - Dynamic linking
  - Engineered to be secured from the ground



# We've come a long way!

- Operating systems were at the core of the revolution
- The **problems** uncovered and the **principles** developed remain fundamental
  - And applicable to many other contexts

# OS History

- *The UNIX Operating System*, AT&T Archives, at YouTube
- *The UNIX Time-Sharing System*, Dennis M. Ritchie and Ken L. Thompson, Bell System Technical Journal 57(6)
- *The Evolution of the Unix Time-sharing System*, Dennis M. Ritchie
- *A Narrative History of BSD*, Kirk McKusick, at YouTube
- *From L3 to seL4: What Have We Learnt in 20 Years of L4 Microkernels?*, Kevin Elphinstone and Gernot Heiser
- *SOSP History Day*, October 4, 2015

## SOSP History Day

October 4, 2015 — Monterey, California, USA

### Introduction

#### **Fifty Years of Operating Systems**

Peter Denning

[FOREWORD](#)

#### **Overview of the Day**

Jeanna Matthews

[ACM & VIDEO](#) — [SLIDES](#)

#### **The Founding of the SOSP Conferences**

Jack Dennis

[ACM & VIDEO](#) — [SLIDES](#)

### 9am - 10:30am

#### **Perspectives on OS Foundations**

Peter Denning

[ABSTRACT](#) — [ACM & VIDEO](#) — [SLIDES](#)

#### **Perspectives on Protection and Security**

Butler Lampson

[ACM & VIDEO](#) — [SLIDES](#)

#### **Perspectives on System Languages and Abstraction**

Barbara Liskov

[ACM & VIDEO](#) — [SLIDES](#)

### 11am - 12:00pm

#### **Evolution of File and Memory Management**

Mahadev Satyanarayanan (Satya)

[ABSTRACT](#) — [ACM & VIDEO](#) — [SLIDES](#)

#### **Reflections on the History of Operating Systems**



# Class operation

# Staff

- Instructor: Pedro Fonseca, [pfonseca@purdue.edu](mailto:pfonseca@purdue.edu)
  - MPI-SWS -> UW -> Purdue
  - Research interests: Operating systems, hypervisors, distributed systems
- TA: Basavesh Shivakumar, [bammanag@purdue.edu](mailto:bammanag@purdue.edu)
  - Research: Operating systems and concurrency

# Acknowledgements

- Course slides heavily based on content by:
  - Douglas Comer
  - Dongyan Xu
  - Byoungyoung Lee

# Resources

- Online resources:
  - Web page: Anchor
  - Piazza: Online discussion, slides, announcements
  - Blackboard: Grades
- PSOs
- Office hours
- Ask questions

# Textbook

- Required textbook:
  - **Operating System Design: The XINU Approach**, Second Edition, Feb 18, 2015, by Douglas Comer
- Other recommended resources:
  - **Operating Systems: Principles and Practice**, 2nd Edition, by Thomas Anderson and Michael Dahlin
  - **Operating Systems: Three Easy Pieces**, by Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau



# Class pre-requisites

- Comfortable writing **C code** (e.g., know well how pointers and memory management works)
- Understand basic data structures (lists, arrays, etc.)
- Comfortable dealing with **large code bases** (e.g., 10-50k LOCs)
- Courses:
  - CS250 (Computer Architecture), CS251 (Data Structures), CS252 (Systems Programming), CS240 (C Programming), or equivalent

# What to expect?

- **Code intensive** (lots of time invested understanding the code base, writing your own code and debugging)
- Understand and deal with low-level **architecture details** (e.g, assembly, CPU registers, memory regions, physical memory addresses)
- Lectures will be code intensive as well

# Lecture Format

- Help you understand important and hard OS concepts, design, and implementation
- Lectures do not cover everything
  - Not all questions in exams are from lectures
  - But many (80%?) of them are :)
- Your responsibility
  - Attend lectures
  - Read code, textbook, lecture notes
  - Labs, quizzes, exams
  - Periodically check web page
  - Read/participate in Piazza discussion

# PSOs

- Not mandatory, but attendance is highly encouraged
- Labs will be explained in the PSOs
- Schedule:
  - PSO 1: 9:30am - 11:20am, Wed @ HAAS 257
  - PSO 2: 9:30am - 11:20am, Fri @ HAAS 257
  - **No PSOs in the first two weeks of class**
- Online students: Special instructions

# Grading

- Tentative grading criteria:
  - Midterm exam: 20%
  - Final exam: 25%
  - Labs: 50%
  - Quizzes: 5%

# Exams and quizzes

- Exams and quizzes are close-book, close-note
- Online students:
  - Proctored exams
  - Blackboard quizzes

# Late days for labs

- There is **no partial credit for late assignments**
- Granted three grace days (24-hour periods) that can be used for any laboratory any time during the semester
  - The three days can be applied to a single assignment (e.g., a lab) or one day can be applied to each of three assignments
  - Grace days must be used in **increments of one full day**
- Once your three grace days have been used, **no further exceptions will be made**
- Grace days cannot be used to extend the due date beyond the last day of regular classes

# Grading policy for labs

- Projects will be done **individually**
  - General discussion allowed
  - Must be on your own when coding starts
- Don't copy the one you found online
  - Will check your submissions against online solutions and other (including previous) submissions



# Academic Integrity Policy

- Academic integrity: lifeline of education
- Your projects, quizzes and exams must be your own
  - we have a zero tolerance policy towards cheating of any kind and any student who cheats will get:
    - **Zero grade** for the project/quiz/exam at first offense
    - **F grade for the course** thereafter
- Both the cheater and the person(s) who aided the cheater will be held responsible for the cheating

# Course organization

- 1.OS overview
- 2.Computer architecture overview
- 3.Process scheduling and management
- 4.Process synchronization
- 5.Memory management
- 6.Interrupt processing
- 7.Device drivers
- 8.File system

# Your feedback

- Your feedback is very welcome!
  - The earlier, the better :)
- Feedback about the class organization, lectures, PSOs, etc.
- Might help you and your colleagues

**What is the operating system?**

# What is the operating system?

- The software system that manages hardware resources for users and their applications
- Many other definitions

# Types of OSs

- OSs may be invisible to “users”
- We’re going to focus on general-purpose operating systems (XINU, Linux, Windows, Android)

# The three roles of (many) OSs

- **Referee**

- Manages resources shared between different applications/users



- **Illusionist**

- Provide an abstraction of physical hardware to simplify application design



- **Glue**

- Provides common services to facilitate sharing between different applications



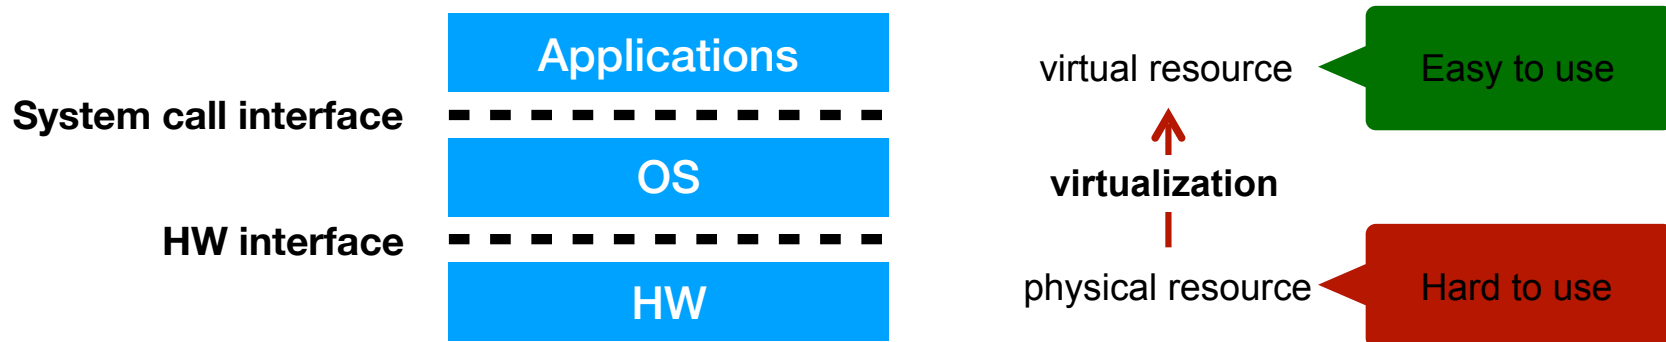
# OS: The referee

- Sharing is central in most system uses
- Sharing raises critical challenges:
  - Resource allocation
  - Isolation
  - Communication
- OS: balances needs, separates conflicts, facilitates sharing



# OS: The illusionist

- *Virtualization* provides the illusion of resources that are not physically present
- A few examples of challenges:
  - HW provides only a small amount of RAM
  - Different HW works in slightly different ways
- OS: virtualizes resources to simplify application development



# OS: The glue

- How to make applications easily share information?
  - OS: provides the communication primitives
- E.g., files, messages, and copy&paste
- A minor role compared with the other two roles
  - Most of the OS complexity arises

# Other systems with similar challenges and designs

- Cloud computing
- Web browsers
- Media players
- Multi-player games
- Multi-user databases
- The internet

**Questions?**