

OS organization

CS503: Operating systems, Spring 2019

Pedro Fonseca
Department of Computer Science
Purdue University

Last lecture

- C and systems resources
 - Books, websites
 - Pre-requisites
 - Typical mistakes

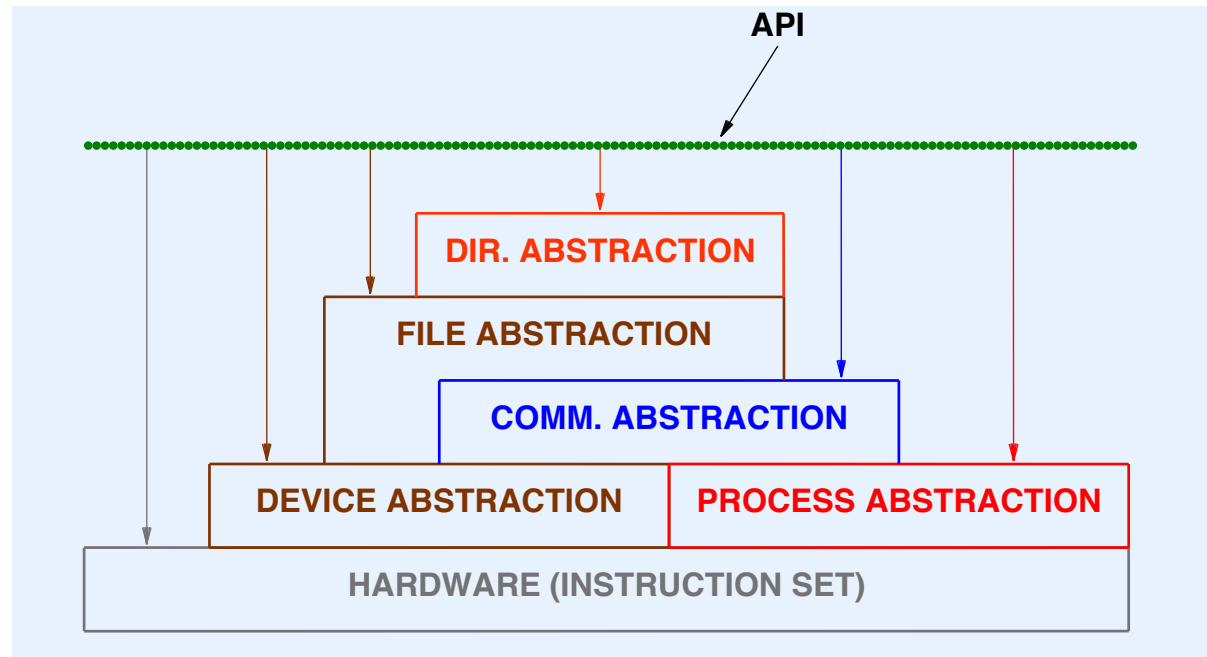
Last lecture

- OS interface:
 - Kernel/user mode
 - Booting process
 - Interrupts
 - System call implementation
 - OS services enabled by HW functions

The system interface

- Single copy of OS per computer
 - Hidden from users
 - Accessible only to application programs
- Application Program Interface (API)
 - Defines services the OS makes available
 - Defines parameters for those services
 - Provides access to all abstractions
 - Hides hardware details

OS Abstractions and Application Interface



- Modules in the OS offer services
- Some services build on others

An OS from outside

Example: System Call in Xinu

- Write a character on console

```
// ex1.c – main
// Write "hi" on the console.

#include <xinu.h>

void main(void) {
    putc(CONSOLE, 'h');
    putc(CONSOLE, 'i');
    putc(CONSOLE, '\n');
    ...
}
```

- We will discuss the implementation of **putc** later

OS services and system calls

- Each OS service accessed through system call interface
- Most services employ a set of several system calls
- Examples:
 - Process management services includes functions to **suspend** and then **resume** a process
 - Socket API used for Internet communications includes many functions

System calls used with I/O

- Open-close-read-write paradigm
- Application:
 - Uses **open** to connect to a file or device
 - Calls functions to **write** data and **read** data
 - Calls **close** to terminate the use
- Internally, the set of I/O functions coordinate:
 - **Open** returns a descriptor **d**
 - **Read** and **write** operate on a descriptor **d**
- Q: Why do we need a descriptor?

Concurrent processing

- Fundamental concept that dominates OS design
- *Real concurrency* achieved by hardware:
 - I/O device operates at the same time as processor
 - Multiple processors/cores each operate at the same time
- *Apparent concurrency* achieved with multitasking (multiprogramming)
 - Multiple programs appear to operate simultaneously
 - Operating system provides the illusion

Multitasking

- Powerful abstraction
- Allows user(s) to run multiple computations
- OS switches processors(s) among available computations quickly
- All computations appear to proceed in parallel

Terminology used with multitasking

- *Program* consists of static code and data
- *Function* is a unit of application program code
- *Process* (also called thread of execution) is an active computation:
 - The execution or “running” of a program

Process

- OS abstraction
- Create by OS system call
- Managed entirely by OS; unknown to hardware
- Operates (i.e., runs) concurrently with other processes

Example of process creation in XINU

```
// ex2.c
// Example of creating processes in Xinu

#include <xinu.h>

void main(void) {
    resume( create(sndA, 1024, 20, "process 1", 0) );
    resume( create(sndB, 1024, 20, "process 2", 0) );
}

// Repeatedly emit 'A' on the console without terminating
void sndA(void) {
    while(1) putc(CONSOLE, 'A');
}

void sndB(void) {
    while(1) putc(CONSOLE, 'B');
}

// Fixed: Lecture slide
// had a typo on the name
```

Difference between function call and process creation

- Function call:
 - Synchronous execution
- Process creation (i.e., `create()` syscall):
 - Asynchronous execution
 - Two processes are running after create
 - ...

Distinction between Program and Process

- Program:
 - Set of functions (code)
 - And initial data
- Process:
 - Computational abstraction
 - Not usually part of the programming language
 - Key idea: Multiple processes can execute the same program

Storage allocation when multiple processes execute

- Various models exist for multitasking environments
- Each process requires its own:
 - Runtime stack
 - Local variables
 - Copy of arguments

Example: Two processes sharing code

```
// Example of 2 processes executing the same code concurrently

#include <xinu.h>

void main(void) {
    resume( create(sndch, 1024, 20, "send A", 1, 'A') );
    resume( create(sndch, 1024, 20, "send B", 1, 'B') );
}

// Output a character on a serial device indefinitely
char sndch(char ch) {
    while (1)      putc(CONSOLE, ch);
}
```

An OS from the inside

How to build an OS

- Work one level at a time
- Identify a service to be provided
- Begin with a **philosophy**
- Establish **policies** that follow the philosophy
- Design **mechanisms** that enforce the policies
- Construct an **implementation** for specific hardware

Design example

- Example: access to IO
- Philosophy: fairness
- Policy: FCFS resource access
- Mechanism: queue of requests (FIFO)
- Implementation: program written in C

Lists and queues

- Fundamental data structures for OS
- Various forms:
 - FIFOs
 - Priority lists
- Operation:
 - Insert item
 - Extract item
 - Delete arbitrary item

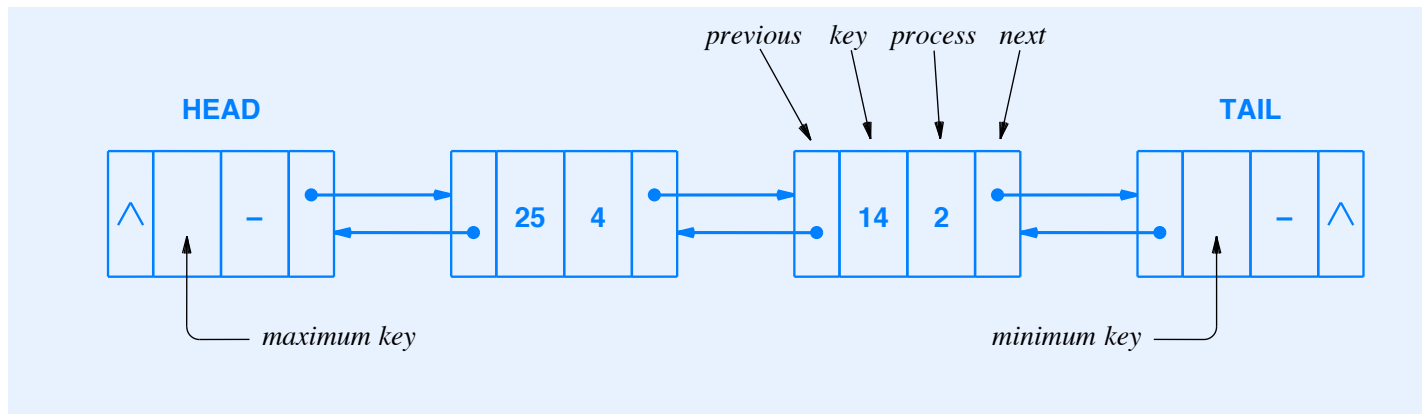
Lists and queues in XINU

- Important ideas:
 - Many lists store processes
 - A process is known by an integer process ID
 - A list store set of process IDs
- A single data structure can be used to store many types of lists

Unified list storage in XINU

- Basic backing storage: **queuetab**
 - A single array used to implement doubly-linked lists
 - Array index represents **pid**
 - Each array element holds: **prev**, **next** and **key**
 - Conceptually each node also holds the process id
- Versatile uses:
 - FIFO (key is not used)
 - Priority list (key is the priority)
- Head and tail are stored elements store in a reserved space in queuetab:
 - Head holds the maximum integer as a key
 - Tail holds the minimum integer as a key

Unified list storage in XINU



- Conceptually nodes contain: **prev**, **next**, **pid**, **key**
- Each list has a **head** and **tail**
- Key is optional (list dependent)

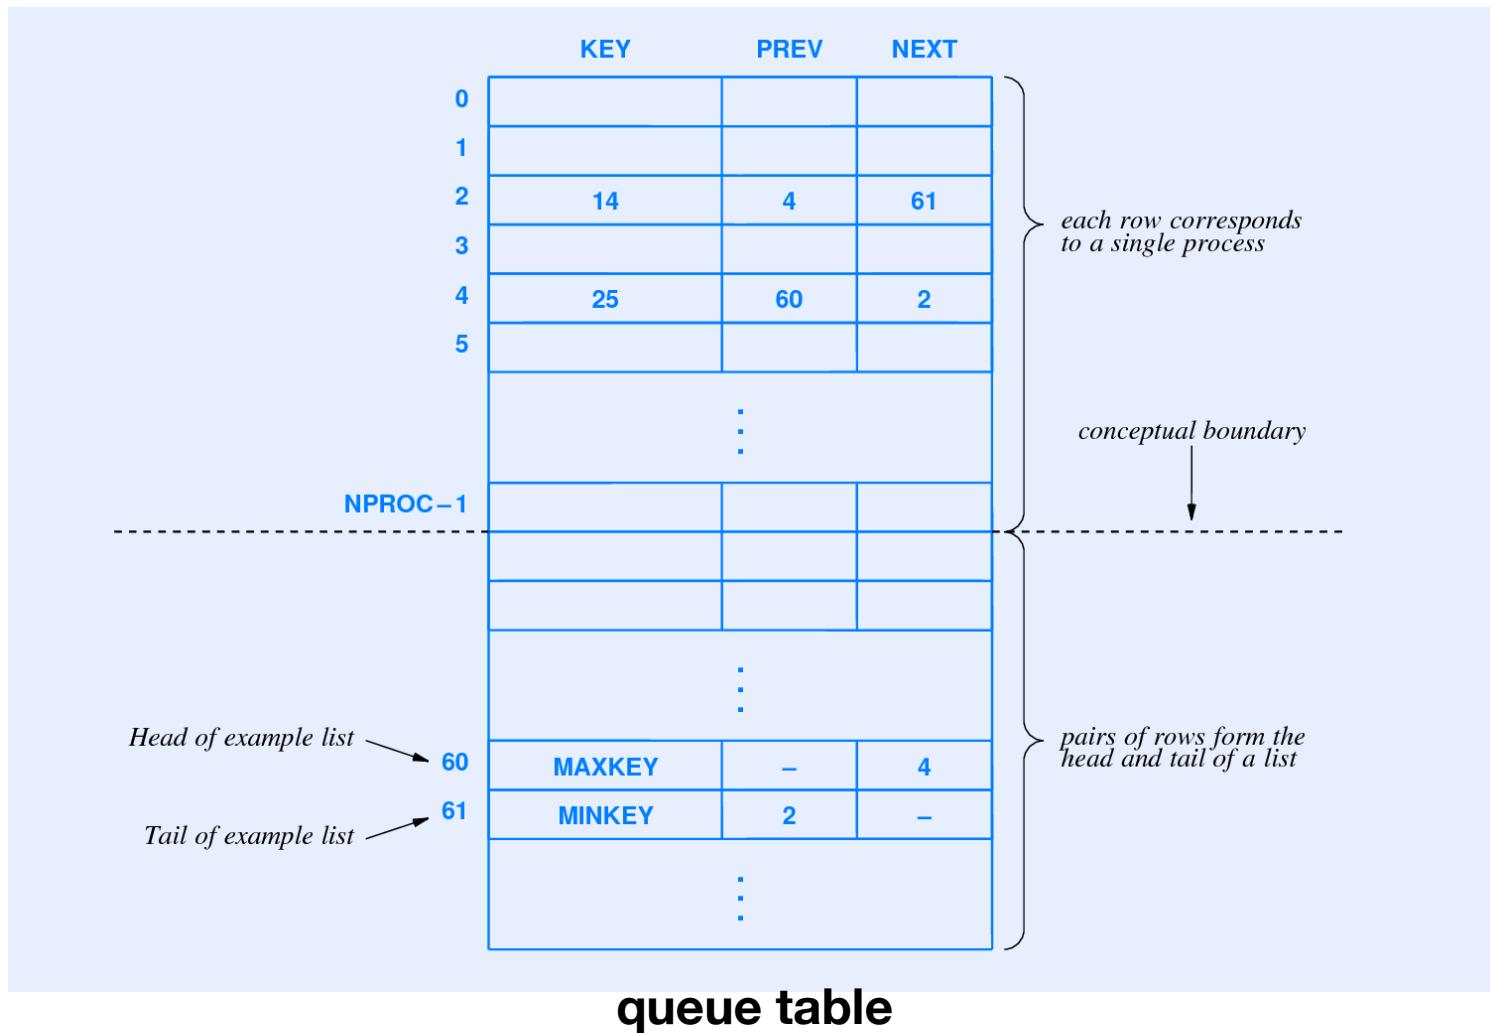
Empty list?

- Two nodes: head and tail node

XINU lists

- Important: Each process can only be in one list
 - This allows some simplifications in the implementation
- Number of processes and list is bounded and small
 - And known during compile time
- Memory usage is important
 - Specially for embedded systems

Unified list storage in XINU



A question about types in C

- K&R C defines **short**, **int** and **long** to be machine-dependent
- ANSI C leaves **int** as a machine-dependent type
- A programmer can define type names
- Q: Should a type specify
 - The purpose of an item?
 - The size of an item?
- Example: Should a process ID type be named:
 - `processid_t` to indicate the purpose?
 - `int32` to indicate the size?

Type names used in Xinu

- Xinu types encompass both purpose and size
- Example: consider a variable that holds an index of the queuetab
 - The type name can specify:
 - That the variable is a queue table index
 - That the variable is a 16-bit signed integer
 - Xinu uses the type name qid16 to specify both

Summary

- Operating systems supply a set of services
- System calls provide interface between OS and application
- Concurrency is a fundamental concept
 - Between IO devices and processor
 - Between multiple computations
- Process an OS abstraction for concurrency
- Process differs from program and function
- Queues and lists