

Process management

CS503: Operating systems, Spring 2019

Pedro Fonseca
Department of Computer Science
Purdue University

Admin

- Labs due next Monday

Previous lecture

- Memory system
 - Endianness, caches
- Bus operations
- Calling conventions
- I/O communication
 - Interrupts
 - DMA

Previous lecture

- DMA:
 - Advantages:
 - Large transfers are more efficient
 - CPU and device *can* work in parallel
 - Disadvantages:
 - Higher per transfer overhead:
 - Need to configure the DMA controller
 - Only compensates for large transfers
 - Cache coherence problems

Goals for today



- How to manage concurrent executions in an operating system?
- What structures the OS needs to keep track off?
- What is and how is a context switch implemented?

Terminology

- The term **process management** has been used for decades to encompass the part of an operating system that manages concurrent execution, including both processes and the threads within them
- The term **thread management** is newer, but sometimes leads to confusion because it appears to exclude processes

Revisit: concurrent processing

- Unit of computation
- Abstraction of a processor
 - Known only to operating systems
 - Not known by hardware

A fundamental principle

- All computations must be done by a process
 - No execution by the OS itself
 - No execution outside of a process
- Key consequence:
 - At any time a process (or a thread) must be running
 - Operating system cannot stop running a process unless it switches to another process

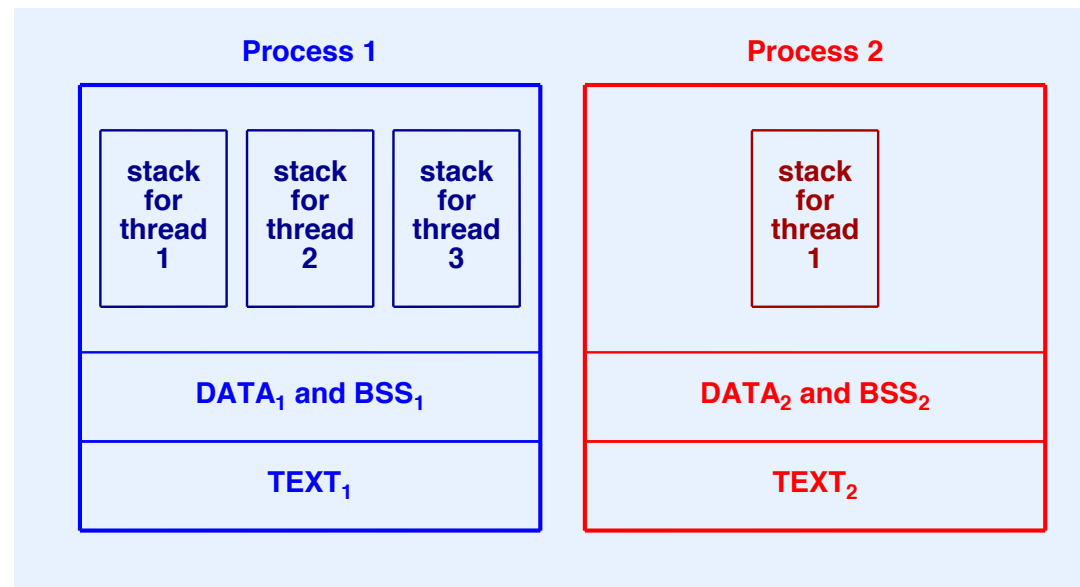
Heavyweight Process (or Process)

- Pioneered in Mach and adopted by Linux
- Address space in which multiple threads can execute
- One data segment per process
- One BSS segment per process
- Multiple threads per process

Lightweight Process (or Thread)

- Single “execution”
- Can share data (data and BSS segments) with other threads
- Must have a private stack segment
 - To keep track of execution contexts
 - Storing local variables, call stacks, etc.

Illustration of two heavyweight processes and their threads



- Threads within a process share text, data, and bss
- No sharing between Processes
- Threads within a Process cannot share stack

Maintaining processes

- Process
 - OS abstraction
 - Unknown to hardware
 - Created dynamically
- Information kept by OS
- OS stores information in a central data structure
 - Called process table
 - Part of OS address space

Information kept in a process table

- For each process:
 - Unique process identifier
 - Owner (a user)
 - Scheduling priority
 - Location of code and data (stack)
 - Process status (running, ready, wait, etc.)
- Q: How about **current program counter** and **current register** values?

Information kept in a process table

- For each thread (if used):
 - Owning process
 - Thread's scheduling priority
 - Location of stack
 - Status of computation

XINU process model

- Simple(st) scheme
- Single-user system (no ownership)
- One global context
- One global address space
- No boundary between OS and applications
- Note: all Xinu processes can share data

Xinu's process table

- **prstate**: The current status of the process (e.g., whether the prstate process is currently executing or waiting)
- **prprio**: The scheduling priority of the process
- **prstkptr**: The saved value of the process stack pointer when the process is not executing
- **prstkbase**: The address of the base of the proces's stack
- **prstklen**: A limit on the maximum size that the process's stack can grow
- **prname**: A name assigned to the process that humans use to identify the process's purpose

Process states

- Used by OS to manage processes
- Set by OS whenever process changes status (e.g., waits for IO)
- Small enum value (integer) stored in the process table
- Tested by OS to determine:
 - Whether a requested operation is valid
 - The meaning of an operation

Process states

- Specified by OS designer
- One “state” assigned per activity
- Values are updated in a process table when activity changes
- Example values:
 - **Current:** process is currently executing
 - **Ready:** process is ready to execute
 - **Waiting:** process is waiting on semaphore
 - **Receiving:** process is waiting to receive a message
 - **Sleeping:** process is delayed for specified time
 - **Suspended:** process is not permitted to execute

Definition of Xinu process state constants

```
// Process state constants
```

```
#define PR_FREE      0    /* Process table entry is unused    */
#define PR_CURR      1    /* Process is currently running     */
#define PR_READY     2    /* Process is on ready queue        */
#define PR_RECV      3    /* Process waiting for message      */
#define PR_SLEEP     4    /* Process is sleeping              */
#define PR_SUSP      5    /* Process is suspended             */
#define PR_WAIT      6    /* Process is on semaphore queue    */
#define PR_RECTIM    7    /* Process is receiving with timeout */
*/
```

- States are defined as needed when system is constructed
- We will understand the process of each state as we consider the system design

Scheduling and context switching

Scheduling

- Fundamental part of process management
- Performed by the OS
- Three steps:
 - Examine the processes that are eligible for execution
 - Select a process to run
 - Switch the processor to the selected process

Implementation of scheduling

- The OS needs a **scheduling policy** that specifies which process to select
- We must then build a scheduling function that:
 - Selects a process according to the policy
 - Update the process table for the current and selected processes
 - Perform a **context switch**:
 - Switching from current process to the selected process

Scheduling policy

- Determines which process is selected for execution
- Typically one of the goals is fairness
- May depend on:
 - User
 - How many processes a user owns
 - Time a given process has been waiting to run
 - Priority of the process
- Note: both hierarchical or flat scheduling can be used

Example of scheduling policy in Xinu

- Each process assigned a priority:
 - Non-negative integer value
 - Initialized when process created
 - Can be changed at any time
- Scheduler always chooses to run an eligible process that has highest priority
- Policy is implemented by a system-wide invariant

The Xinu scheduling invariant

- “At any time, the processor must be executing a highest priority eligible process. Among processes with equal priority, scheduling is round robin.”
- Invariant must be enforced whenever:
 - The set of eligible processes changes
 - The priority of any eligible process changes
- Such changes only happen during:
 - A system call
 - An interrupt

Implementation of scheduling

- Process is eligible if state is **ready** or **current**
- To avoid searching process table during scheduling
 - Keep ready processes on linked list called **ready list**
 - Order the ready list by process priority
 - Selection of highest-priority process can be performed in a constant time (i.e., $O(1)$)

High-speed scheduling decision

- Compare priority of **current** process to priority of first process on **ready list**
 - If current process has a higher priority, do nothing
 - Otherwise, extract the first process from the ready list and perform a context switch to switch the processor to the next process

Deferred rescheduling

- Delays enforcement of scheduling invariant
- Prevents rescheduling temporarily
 - A call to **resched_ctrl(DEFER_START)** suspends rescheduling
 - A call to **resched_ctrl(DEFER_STOP)** resumes normal scheduling
- Main purpose: allow device driver to make multiple processes ready before allowing any of them to run
- We will see an example later

Xinu scheduler details

- Before calling the scheduler:
 - Global variable **currpid** gives ID of process that is executing
 - **proctab[currpid].prstate** must be set to the desired next state for the current process
 - Q: How do you reference **currpid** if running on multi-core CPUs?
- If current process remains eligible and has highest priority, scheduler does nothing
 - Just returns and keeps executing the current process
- Otherwise, scheduler moves current process to the specified state and runs the highest priority **ready** process

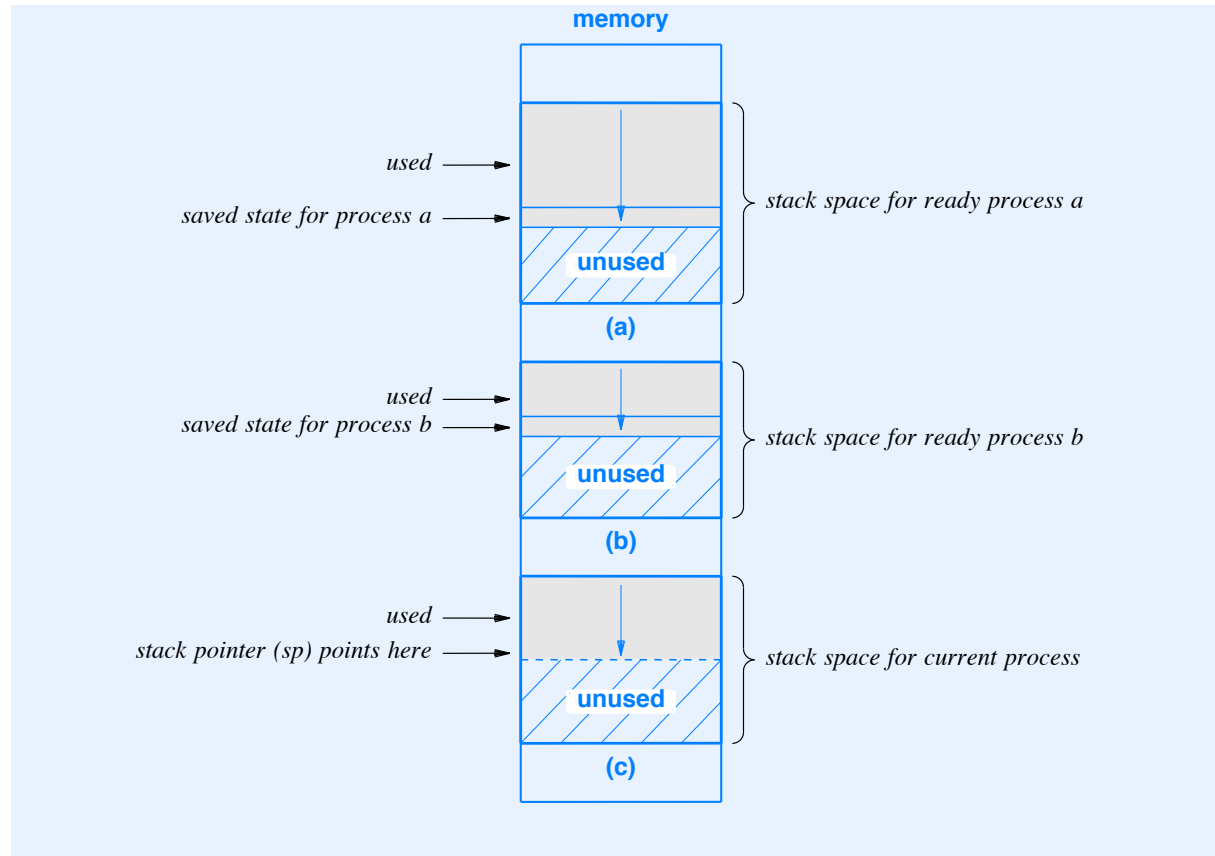
Xinu scheduler details

- What if several processes have the same priority and are the highest priority processes?
- When inserting a process on the ready list, places the process behind other processes with the same priority
 - Round-robin scheduling
- If scheduler switches context, first process on ready list is selected
- Note: scheduler switches context if the first process on the ready list has priority **equal** to the current process

Example scheduler code

- **resched()** and **resched_cntl()**
 - in “**system/resched.c**”

Illustration of states saved on process stack




- Stack of each process contains the saved state

Process state transition

- Recall each process has a state
- State determines:
 - Whether an operation is valid
 - Semantics of each operation
- Transition diagram documents valid operations

Example: Transition between current and ready state



- **resched()** can move a process in either direction

Context switch

- Component of process manager
- Low level (must manipulate underlying hardware)
- Usually written in assembly language
- Called by scheduler
- Moves processor from one process to another

Context switch operation

- Given a new process **N** and old process **O**:
- Save copy of all information pertinent to O on process O's stack or the process table entry for process O
 - Contents of HW registers
 - PC
 - Privilege level and hardware status
 - Memory map and address space
- Load saved information for N
- Resume execution of N

Context switch stack

```
+-----+
|...    | <- ESP
|regs   | (pushed by pushal, popped by popal)
|...    |
+-----+
|flag reg| (pushed by pushfl, popped by popfl)
+-----+
|SFP     | <- EBP
+-----+
|RET-ADDR| (return address of ctxsw)
+-----+
|old_sp  | <- EBP+8
+-----+
|new_sp  | <- EBP+12
+-----+
```

low address values

high address values

(note diagram is in inverted order)

Context switch code (fixed)

- `ctxsw()` in `system/ctxsw.S`

```
10 ctxsw:
11     pushl    %ebp                /* Push ebp onto stack */
12     movl     %esp,%ebp          /* Record current SP in ebp */
13     pushfl                   /* Push flags onto the stack */
14     pushal                   /* Push general regs. on stack */
15
16     /* Save old segment registers here, if multiple allowed */
17
18     movl     8(%ebp),%eax        /* Get mem location in which to */
19                                     /* save the old process's SP */
20     movl     %esp,(%eax)        /* Save old process's SP */
21     movl     12(%ebp),%eax      /* Get location from which to */
22                                     /* restore new process's SP */
23
24     /* The next instruction switches from the old process's */
25     /* stack to the new process's stack. */
26
27     movl     (%eax),%esp        /* Pop up new process's SP */
28
29     /* Restore new seg. registers here, if multiple allowed */
30
31     popal                   /* Restore general registers */
32     movl     4(%esp),%ebp        /* Pick up ebp before restoring */
33                                     /* interrupts */
34     popfl                   /* Restore interrupt mask */
35     add      $4,%esp            /* Skip saved value of ebp */
36     ret                          /* Return to new process */
```

Book and lab code also include two important instructions

Summary

- Process management
- Process structures
- Scheduling and scheduling policies
- Context switch