

Process management 2

CS503: Operating systems, Spring 2019

Pedro Fonseca
Department of Computer Science
Purdue University

Admin

- One more TA: Hafiz Kamran Khalil :)
- Next lab is scheduled to be released next week

Previous lecture

- Process abstraction
 - Heavy process vs. light process
- Process structures kept by the kernel
 - Process table
- Scheduling and scheduling policies
- Context switch
 - $O(1)$ only for picking the first process
 - But insertion takes more time see the implementation

Goals for today



- Continue understanding process management:
 - How to manage concurrent executions in an operating system?
 - What structures the OS needs to keep track off?
 - What is and how is a context switch implemented?
- What are the main process management services?

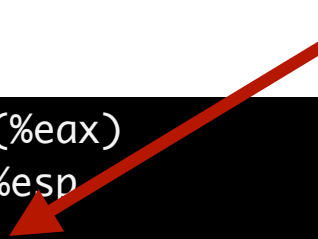
How are arguments passed in Xinu?

- `Objdump -S xinu.elf`

How are arguments passed in Xinu?

Arg

```
113dba:    c7 00 00 00 00 00    movl    $0x0, (%eax)
113dc0:    83 ec 0c              sub     $0xc, %esp
113dc3:    6a 02                push    $0x2
113dc5:    e8 1a 0d ff ff       call    104ae4 <resched_cntl>
113dca:    83 c4 10              add     $0x10, %esp
113dcd:    83 ec 0c              sub     $0xc, %esp
113dd0:    ff 75 f4              pushl   -0xc(%ebp)
```



```
51 /*-----
52 *  resched_cntl  -  Control whether rescheduling is deferred or allowed
53 *-----
54 */
55 status  resched_cntl(          /* Assumes interrupts are disabled */
56         int32 defer           /* Either DEFER_START or DEFER_STOP */
57         )
58 {
```

Assembly background

- Assembly:
 - one statement -> one machine instruction
 - assembler directives and labels
 - instruction format is restricted by architecture
- AT&T assembly syntax format:
 - [instruction mnemonic] [src] [destination]
 - Intel manuals use another format (i.e., Intel assembly syntax)
- Operands:
 - Register value: *%rax*
 - Memory location pointed by a register value: *(%rax)*
 - Memory location of a register value + offset: *12(%ebp)*
 - Some instructions change implicit registers (e.g., flags is changed by CMP and SUB)
 - Some instructions read implicit registers (e.g., JMP reads flags)

More info: https://en.wikibooks.org/wiki/X86_Assembly/GAS_Syntax

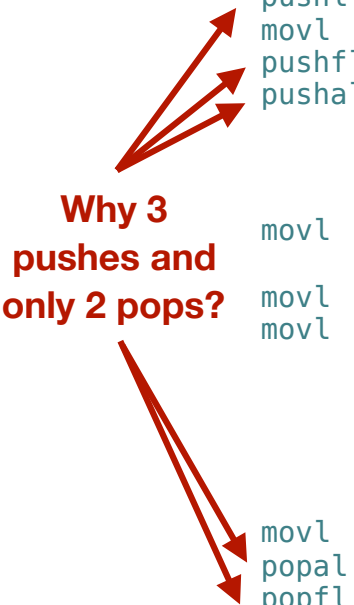
Context switch

- Restoring the CPU changes the CPU state itself
- Challenge is to make it atomic and avoid loss of information
- Particularly challenging is the PC register
 - It cannot even be read explicitly in x86

Context switch code

- `ctxsw()` in `system/ctxsw.S`

```
/*-----  
 * ctxsw - X86 context switch; the call is ctxsw(&old_sp, &new_sp)  
 *-----  
 */  
ctxsw:  
    pushl    %ebp          /* Push ebp onto stack      */  
    movl     %esp,%ebp     /* Record current SP in ebp */  
    pushfl   %eax          /* Push flags onto the stack */  
    pushal   %eax          /* Push general regs. on stack */  
  
                                /* Save old segment registers here, if multiple allowed */  
    movl     8(%ebp),%eax   /* Get mem location in which to */  
                                /*   save the old process's SP   */  
    movl     %esp,%eax     /* Save old process's SP      */  
    movl     12(%ebp),%eax /* Get location from which to */  
                                /*   restore new process's SP   */  
  
                                /* The next instruction switches from the old process's */  
                                /*   stack to the new process's stack.          */  
  
    movl     (%eax),%esp    /* Pop up new process's SP    */  
    popal    %eax          /* Restore general registers  */  
    popfl    %eax          /* Restore interrupt mask     */  
    ret
```



Why 3 pushes and only 2 pops?

(Buggy code in previous lecture)

Context switch code

- `ctxsw()` in `system/ctxsw.S`

```
10 ctxsw:
11     pushl    %ebp                /* Push ebp onto stack */
12     movl     %esp,%ebp          /* Record current SP in ebp */
13     pushfl                   /* Push flags onto the stack */
14     pushal                   /* Push general regs. on stack */
15
16     /* Save old segment registers here, if multiple allowed */
17
18     movl     8(%ebp),%eax        /* Get mem location in which to */
19                                     /* save the old process's SP */
20     movl     %esp,(%eax)        /* Save old process's SP */
21     movl     12(%ebp),%eax      /* Get location from which to */
22                                     /* restore new process's SP */
23
24     /* The next instruction switches from the old process's */
25     /* stack to the new process's stack. */
26
27     movl     (%eax),%esp        /* Pop up new process's SP */
28
29     /* Restore new seg. registers here, if multiple allowed */
30
31     popal                   /* Restore general registers */
32     movl     4(%esp),%ebp        /* Pick up ebp before restoring */
33                                     /* interrupts */
34     popfl                   /* Restore interrupt mask */
35     add      $4,%esp            /* Skip saved value of ebp */
36     ret                          /* Return to new process */
```

Book and lab code also include two important instructions

Context switch

- Restoring the CPU changes the CPU state itself
- Challenge is to make it atomic and avoid loss of information
- Particularly challenging is the PC register
 - It cannot even be read explicitly in x86
 - *The trick is to not store the current PC value; instead we store the location where the `ctxsw()` should return*
 - *Using the **stack** and the **call/ret** instruction semantics*

Evolution of the x86 context switch in Linux

- Many ways to do the context switch in x86!
- Linux when through a long evolution
- HW support for context switching
- Multiprocessor challenges:
 - Migration of processes between CPUs
 - Each processor has a dedicated interrupt controller (APIC)

More info at:

http://www.maizure.org/projects/evolution_x86_context_switch_linux/

Hardware context switching

- Also known as “Hardware Task Switching” in CPU manuals
- May only handle part of the process state (e.g.,. exclude FPU/MMX and SSE state)
- In x86: one mechanism relies on the task register (TR) and a far CALL or JMP to load a task state segment (TSS)
- Hardware context switching (vs. software context switching)
 - Exact mechanisms is very hardware specific
 - May perform unnecessary checks (e.g., segmentation)
 - In practice slow

Question #1

- Our invariant says that at any time, a process must be executing
- Context switch code moves from one process to another
- **Q:** which process executes the context switch code?

Question #2

- At any time, CPU should execute something
- **Q:** If all user processes may be idle (e.g., applications all wait for input), which process should OS executes?

Null process

- Does not compute anything useful
- Is present merely to ensure that at least one process remains ready at all times
- Simplifies scheduling (no special cases)

Code for a null process

- Easiest way to code a null process:

```
while (1);  
    /* Do nothing */
```

- May not be optimal. **Q:** Why?
- Fetch-execute takes bus cycles
- **Q:** How to optimize it?
- Two possible ways:
 - Some processors offer a special pause instruction that stops the processor until an interrupt occurs
 - Instruction cache can eliminate bus accesses

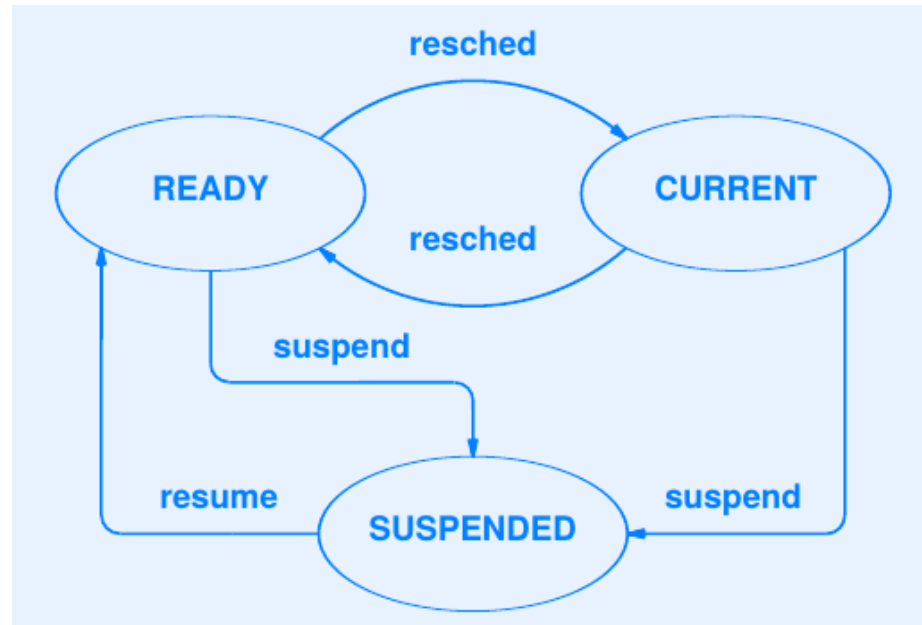
Process manipulation

- Need to invent ways to control processes
- Example operations:
 - Suspension
 - Resumption
 - Creation
 - Termination
- Recall: state variable in process table store the status

Process suspension

- Temporarily “stop” a process
- Prohibit it from using the processor
- To allow later resumption:
 - Process table entry retained
 - Complete state of computation saved
- OS sets process table entry to indicate process is suspended

State transition for suspension and resumption



- Either current or ready process can be suspended
- Only a suspended process can be resumed
- System calls can suspend or resume

Suspension code

- `suspend()` function

Process resumption

- Resume the execution of previously suspended process
- Method:
 - Make process eligible for processor
 - Re-establish scheduling invariant
- Note: resumption does not guarantee instantaneous execution

Resumption code

- `resume()` function

Function to make the process ready

- ready() function
- Make a process eligible for execution

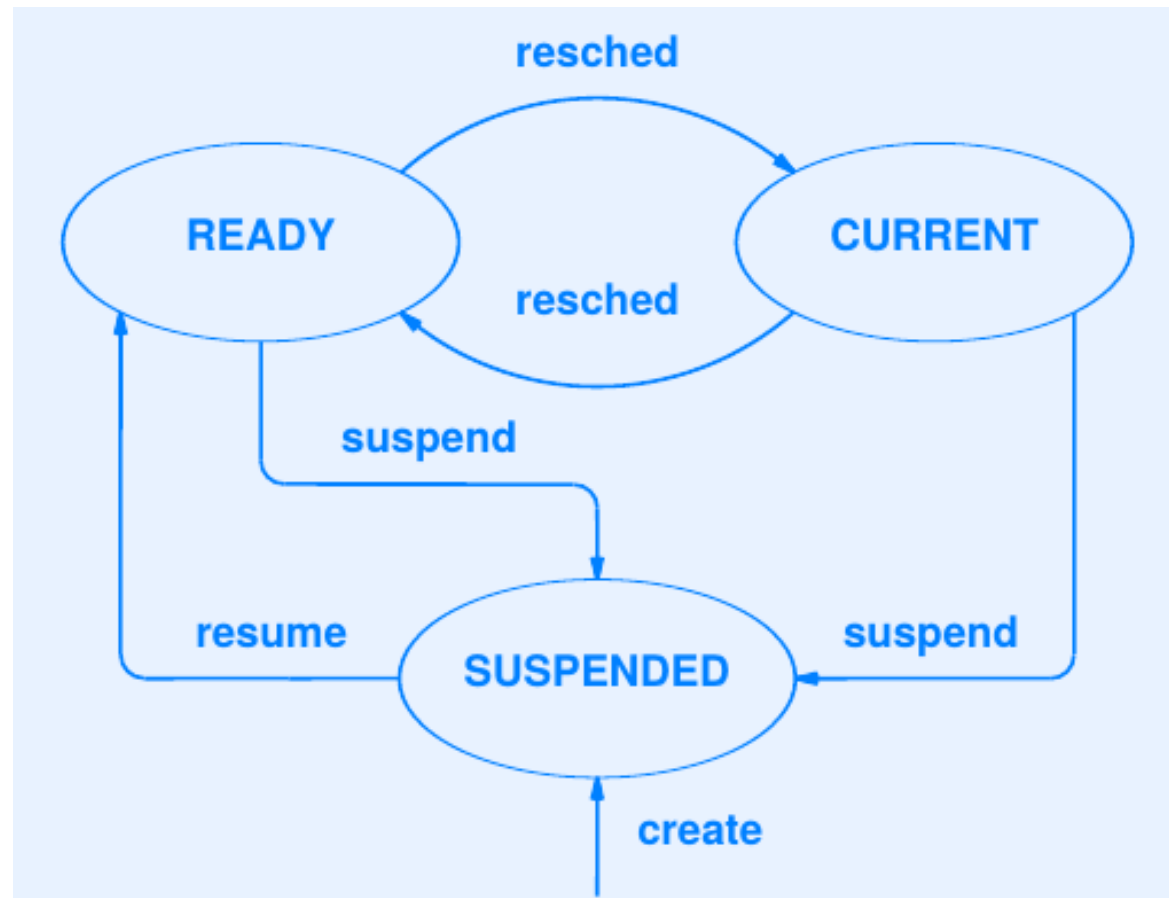
Process termination

- Final and permanent
- Record of a process is removed
- Process table entry becomes available for reuse
- Known as process exit if initiated by the thread itself
- We will see more about termination later

Process creation

- Processes are dynamic — process creation refers to starting a new process
- Implemented by **create** function in Xinu
- Method:
 - Find free entry in process table
 - Fill in entry
 - Place new process in suspended state
- Needs to select a new pid

State transition for additional process management functions



Note that both current and ready processes can be suspended

Summary

- Context switch
- Null process
- Process management functions:
 - Suspension
 - Resumption
 - Creation
 - Termination