

Memory management 4

CS503: Operating systems, Spring 2019

Pedro Fonseca
Department of Computer Science
Purdue University


Previous lecture (s)

- Memory allocation
- Paging
 - Paging divides the memory into pages/page frames
 - Enables mapping logical addresses into physical addresses at the page granularity
 - High-degree of flexibility for OSs
 - Causes more memory accesses
 - Relies on hardware (MMU) and in-memory configuration structure (page table)

Recall: CPU access to memory


- The CPU reads instructions and reads/writes data from/to memory
- Functional interface:

```
- `value = read(address)`  
- `write(address, value)`
```



Recall: Logical addressing

- Memory management unit (MMU)
 - Real-time, on demand translation between logical (virtual) and physical addresses



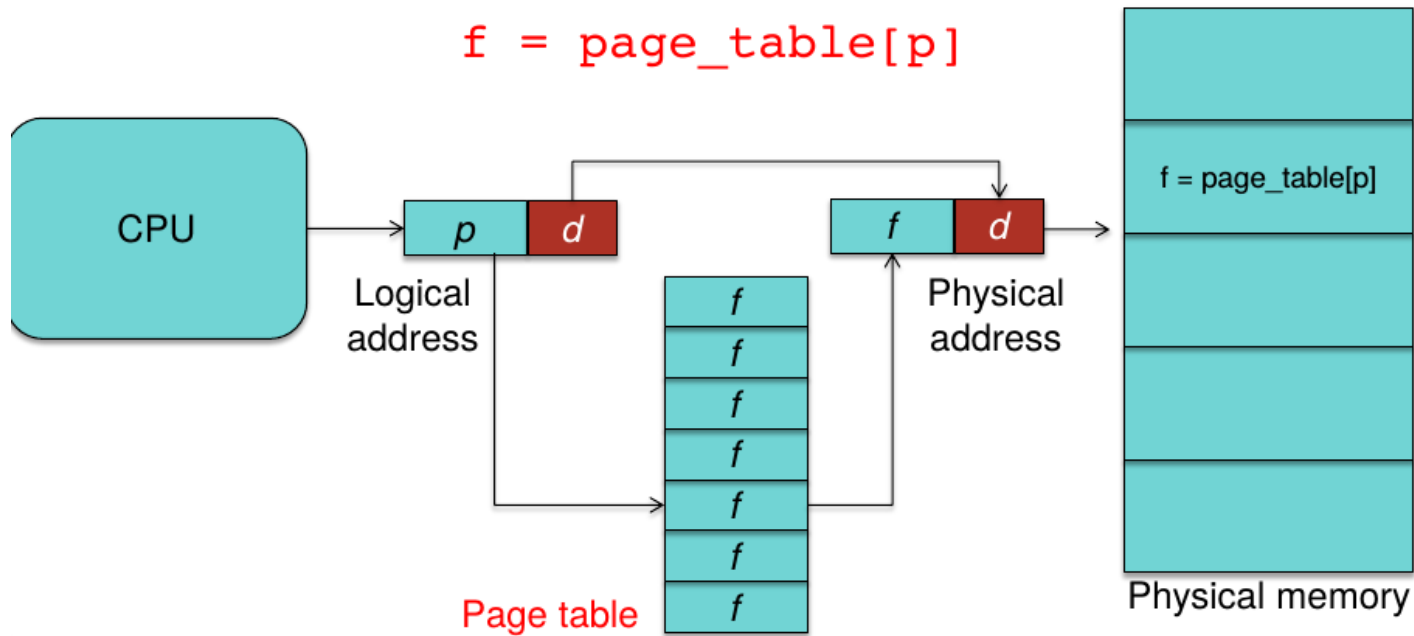
Recall: Paging

- Memory management scheme
 - Physical space can be non-contiguous
 - No fragmentation problems
 - No need for compaction
- Paging is implemented by the MMU (in the CPU)

Recall: Paging

- Translation:
 - Divide physical memory into fixed-size blocks: **page frames**
 - A logical address is divided into blocks of the same size: **pages**
 - All memory accesses are translated:
 - page (virtual) -> page frame (physical)
 - The mapping is stored in a **page table**
- Example:
 - 32-bit address, 4KB page size:
 - Top 20 bits identify the **page number**
 - Bottom 12 bits identify offset with the page or page frame

Recall: Page translation



- p : upper 20-bits in a logical address, page index
- d : lower 12-bits in a logical address, offset within the page or page frame
- f : page frame index

Recall: Logical vs. physical views of memory

Page 3
page 2
page 1
page 0

Logical Memory

3	-	→ page 3 not mapped
2	7	
1	2	
0	4	

Page Table

frame	
7	Page 2
6	
5	
4	Page 0
3	
2	Page 1
1	
0	

Physical Memory

Recall: Hardware-assisted page translation

- Where do you keep the page table?
 - In memory
- How do we find the page table?
 - Page table base register: **CR3** register on x86 architecture
- Software-based translation can be slow!
 - To read a byte of memory, we need to read the page table first!
 - Each memory access would now be x2 (or more) times slower
- Hardware-assisted translation
 - **MMU** and **TLB**

Hardware acceleration: TLB

- Cache mappings for frequently accessed pages:
 - Translation Lookaside Buffer (TLB)
 - Associative memory: key (page #) and value (frame #)
- TLB hit vs miss
 - If there is a TLB miss => need to perform the page table lookup in RAM
- TLB is on-chip and fast, but small (from 64 to 1024 entries)
 - Locality (temporal and spatial) may help to keep good TLB hit rate

Address space identifiers: Tagged TLB

- There is only one TLB per system
 - But there can be multiple address spaces (per process)
- When we context switch, we switch address spaces
 - Some TLB entries may belong to the old address space
- Solutions:
 - Flush / invalidate the entire TLB
 - Have a Tagged TLB with an address space identifier

MMU

- Memory management unit (MMU):
 - A hardware unit (typically on-chip)
 - Walks the page table
- An MMU can also **enforce memory protection**
 - Page table stores status & protection bits per page
 - **valid/invalid**: is there a frame mapped to this page
 - **read-only**:
 - **no-execute**
 - **kernel only access**
 - **dirty**: the page has been modified since the flag was cleared
 - **accessed**: the page has been accessed since the flag was cleared

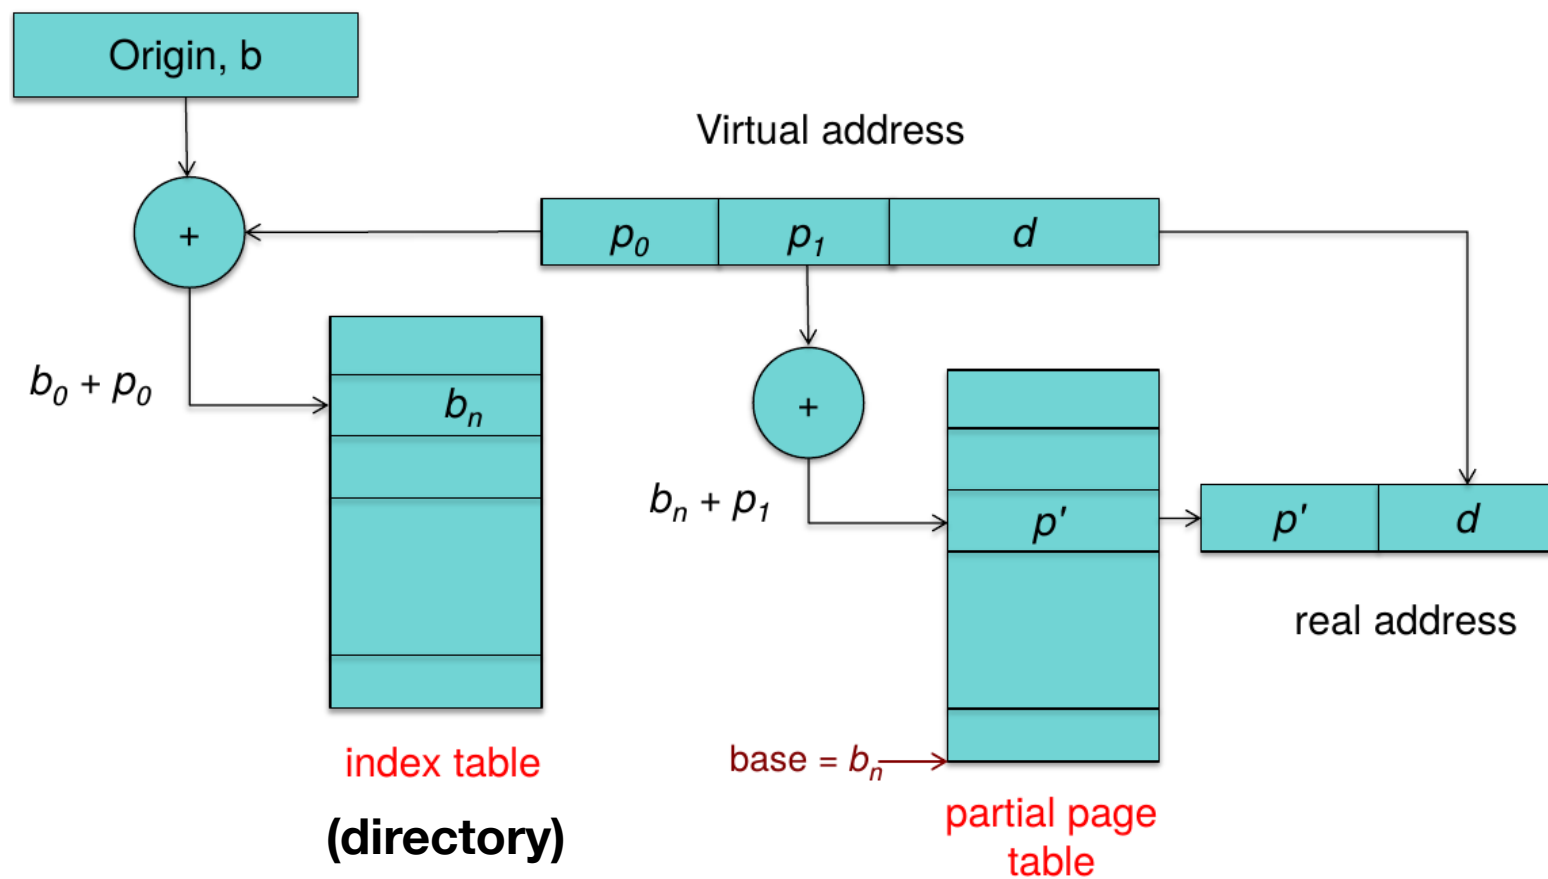
How much space used by the page table?

- Example: 32-bit system with 4KB pages: 20-bit page table
- $\Rightarrow 2^{20} = 1048576$ entries in the page table

Multi-level (hierarchical) page tables

- Most processes use only a small part of their address space
- Keeping a (single-level) page table for the entire address space is wasteful
 - Example:
 - 32-bit system with 4KB pages: 20-bit page table
 - $\Rightarrow 2^{20} = 1048576$ entries in the page table

Multilevel page table example



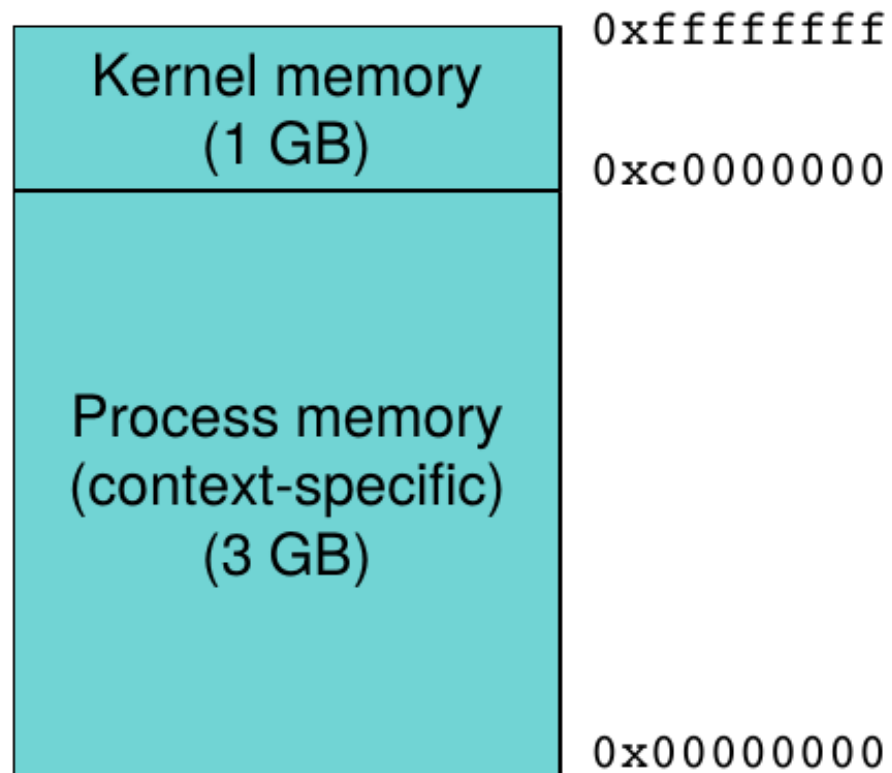
Page-based virtual memory benefits

- Allow discontinuous allocation
 - Simplify memory management for multiprogramming
 - MMU gives the illusion of contiguous allocation of memory
- Process can get memory any where in the address space
 - Allow a process to “feel” that it has more memory than it really has
 - Process can have greater address space than system memory
- Enforce memory protection
 - Each process’s address space is separated from others
 - MMU allows pages to be protected

Typical kernel's view of memory

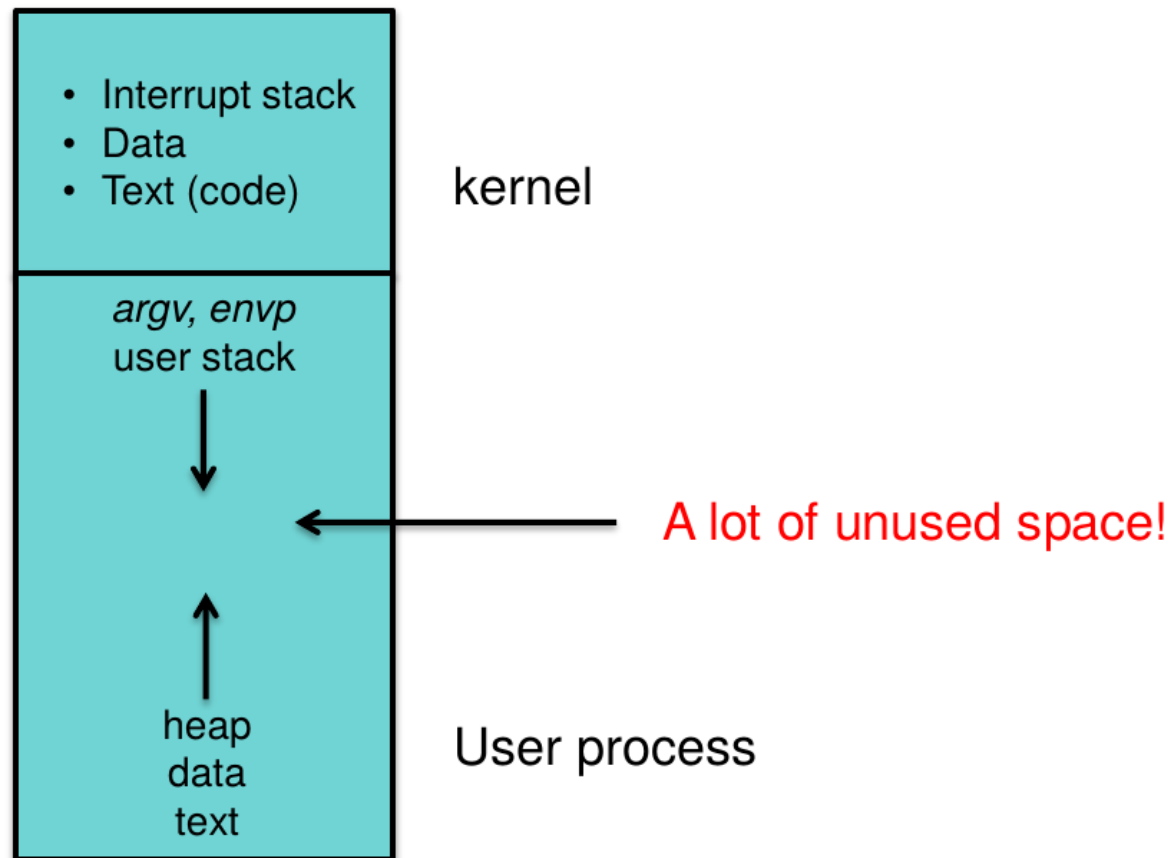
- A process sees a flat linear address space
 - Accessing regions of memory mapped to the kernel causes a page fault
- Kernel's view:
 - Address space is divided into two parts:
 - User part: changes with context switches
 - Kernel part: remains constant across context switches
 - Split is configurable:
 - 32-bit x86: 3GB for user + 1 GB for kernel

Example




Example:

Sample per-process memory map



Virtual memory makes sharing easy



Copy on write

- Share until a page gets modified
- Q: When would this be useful?

Copy on write use case example

- Example: fork()
 - Set all pages to read-only
 - Trap on write
 - If legitimate write:
 - Allocate a new page and copy contents from the original
- Q: Other related use-cases for memory protection?
 - Zero-on-reference

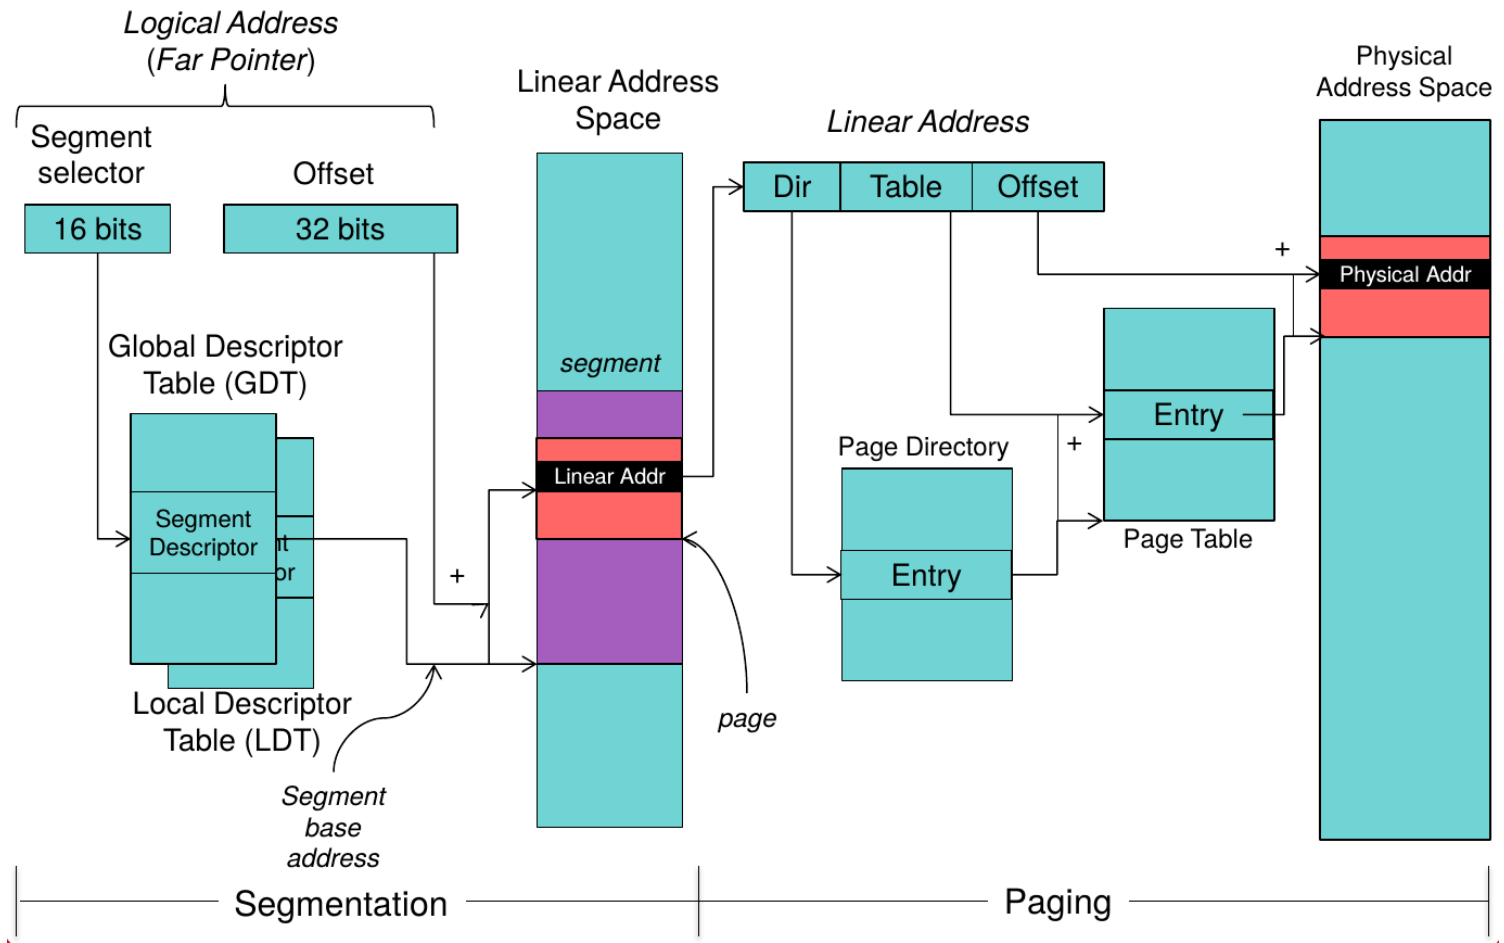
x86 microprocessor modes

- Real mode (IA-32):
 - Memory model: “**Segmentation**”
 - Logical address is expressed as A:B
 - **A**: Segment (value **in** segment register)
 - **B**: Offset within the segment
 - Physical address = $A * 0x10 + B$
 - Segment registers:
 - CS for code, DS for data, SS for stack, ES, FS, and GS
 - Stores the base address of the segment
- Real mode is rarely used nowadays, apart from boot

x86 microprocessor modes

- Protected mode (IA-32):
 - Memory model: **Segmentation + Paging**
 - Logical address is expressed as A:B
 - A: Segment selector (Index of Global Descriptor Table)
 - B: Offset within the segment (32-bits)
 - Logical address = Segment base address (specified by GDT[A]) + B
 - The paging mechanism can be optionally enabled

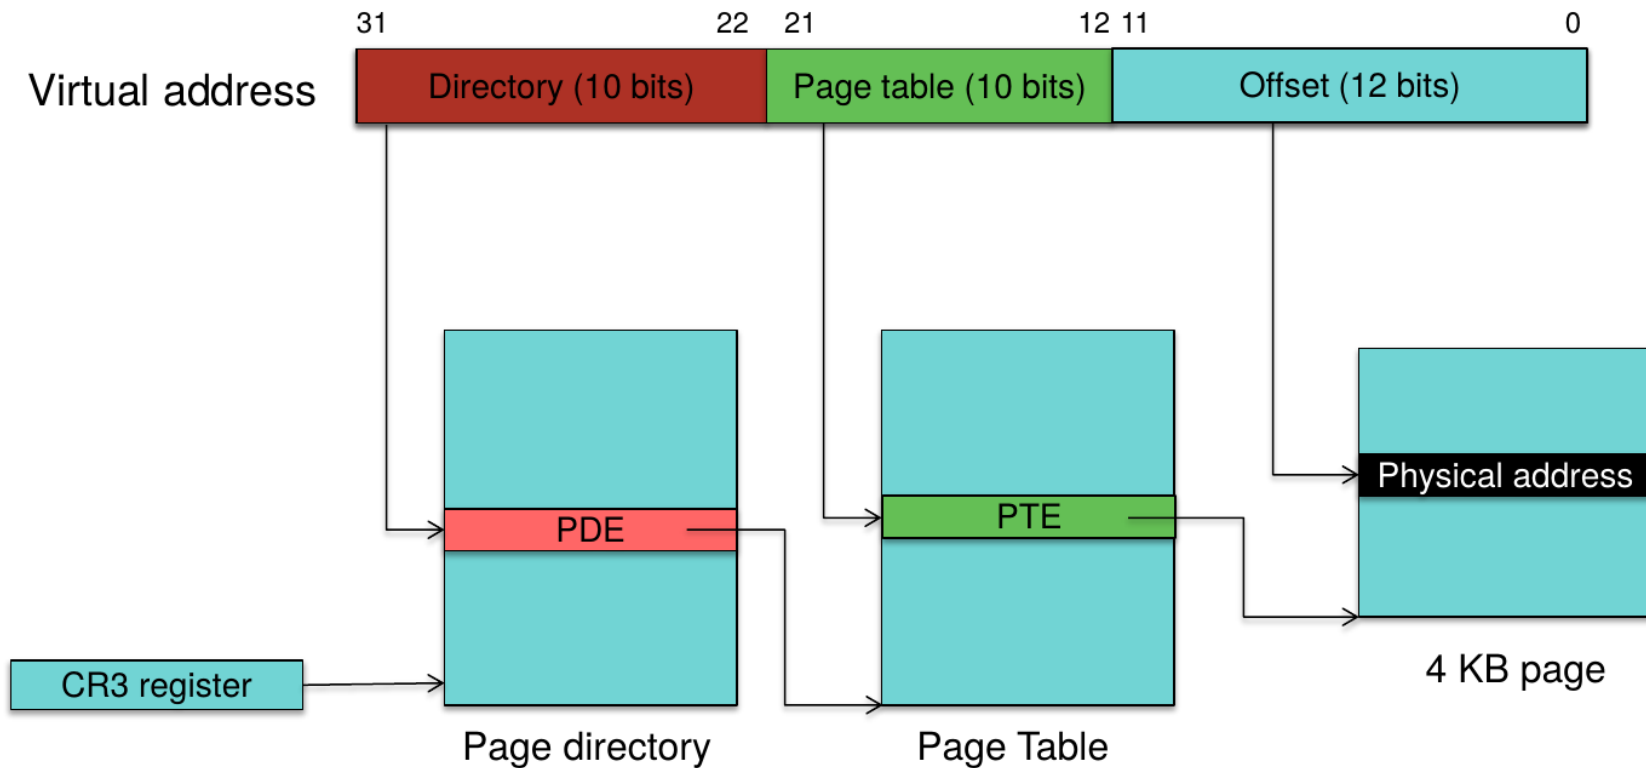
IA-32 segmentation + paging



x86 microprocessor modes

- Long mode (x86-64):
 - The memory model that most of you x86-64 computers run on
 - Segmentation is almost obsolete:
 - Except for FS and GS
 - Use flat memory model (i.e., descriptor base = 0, limit = max)
 - Check more details:
 - <http://wiki.osdev.org/Segmentation>
 - <http://wiki.osdev.org/X86-64>

32-bit paging with 4KB pages



Example: TLBs on the Core i7

- 4 KB pages:
 - Instruction TLB: 128 entries per core
 - Data TLB: 64 entries
- Old CPUs: Cache and TLB specs core 2 Duo

64-byte Prefetching

Data TLB: 4-KB Pages, 4-way set associative, 256 entries

Data TLB: 4-MB Pages, 4-way set associative, 32 entries

Instruction TLB: 2-MB pages, 4-way, 8 entries or 4M pages, 4-way, 4 entries

Instruction TLB: 4-KB Pages, 4-way set associative, 128 entries

L1 Data TLB: 4-KB pages, 4-way set associative, 16 entries

L1 Data TLB: 4-MB pages, 4-way set associative, 16 entries

<http://www.cpu-world.com/sspec/SL/SLB9L.html>