

Memory management 5

CS503: Operating systems, Spring 2019

Pedro Fonseca
Department of Computer Science
Purdue University

Copyright © 2018 by Douglas Comer And CRC Press. All rights reserved.
Edited by Byoungyoung Lee and Pedro Fonseca


Admin

- Midterm on Monday
- Extra-office hour tomorrow:
 - Friday, March 1: 1:30 – 2:30pm @ LWSN 3154N
- Lab2 is due next Friday

Previous lecture

- Paging
 - Processor modes
 - Virtual memory mechanisms:
 - Paging
 - Segmentation
 - Page tables, multi-level paging, PDEs/PTEs, etc.
 - TLB
 - Memory protection, memory layout with virtual memory

Recall: Page translation



- p : upper 20-bits in a logical address, page index
- d : lower 12-bits in a logical address, offset within the page or page frame
- f : page frame index

Recall:

Hardware acceleration: TLB

- Cache mappings for frequently accessed pages:
 - Translation Lookaside Buffer (TLB)
 - Associative memory: key (page #) and value (frame #)
- TLB hit vs miss
 - If there is a TLB miss => need to perform the page table lookup in RAM
- TLB is on-chip and fast, but small (from 64 to 1024 entries)
 - Locality (temporal and spatial) may help to keep good TLB hit rate

Locality in Access Patterns

- **Temporal locality:** recently referenced locations are more likely to be referenced in the near future: e.g., files
- **Spatial locality:** referenced locations tend to be clustered: e.g., listing all files under a directory

Recall:

Address space identifiers: Tagged TLB

- There is only one TLB per system
 - But there can be multiple address spaces (per process)
- When we context switch, we switch address spaces
 - Some TLB entries may belong to the old address space
- Solutions:
 - Flush / invalidate the entire TLB
 - Have a Tagged TLB with an address space identifier

Recall: MMU

- Memory management unit (MMU):
 - A hardware unit (typically on-chip)
 - Walks the page table
- An MMU can also **enforce memory protection**
 - Page table stores status & protection bits per page
 - **valid/invalid:** is there a frame mapped to this page
 - **read-only:**
 - **no-execute**
 - **kernel only access**
 - **dirty:** the page has been modified since the flag was cleared
 - **accessed:** the page has been accessed since the flag was cleared

Recall:

Page-based virtual memory benefits

- Allow discontinuous allocation
 - Simplify memory management for multiprogramming
 - MMU gives the illusion of contiguous allocation of memory
- Process can get memory any where in the address space
 - Allow a process to “feel” that it has more memory than it really has
 - Process can have greater address space than system memory
- Enforce memory protection
 - Each process’s address space is separated from others
 - MMU allows pages to be protected

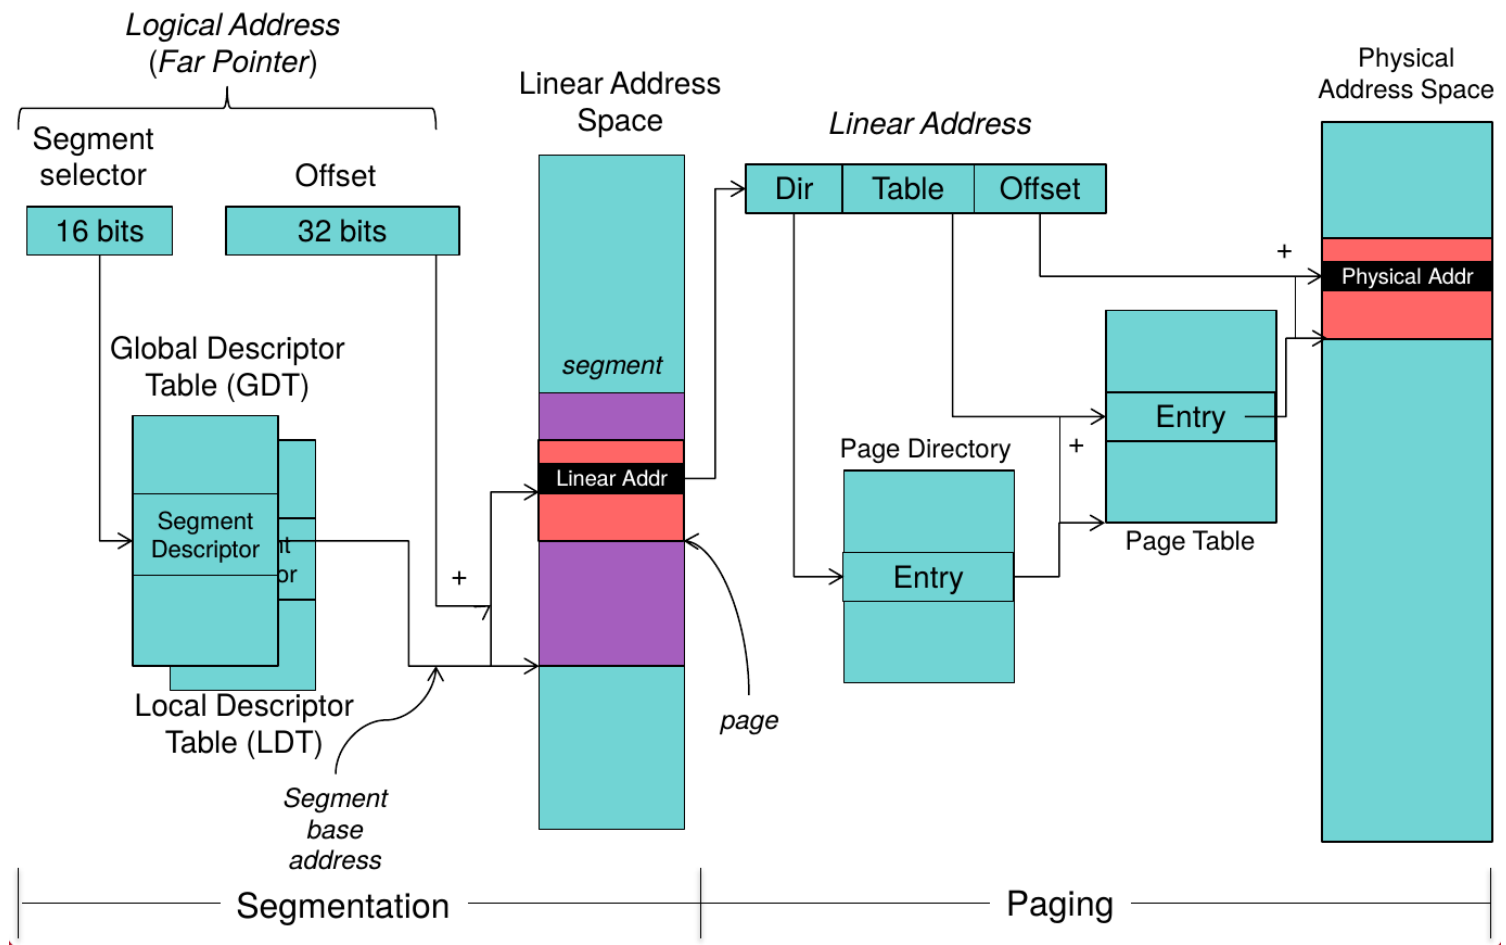
Recall: x86 microprocessor modes

- Real mode (IA-32):
 - Memory model: “**Segmentation**”
 - Logical address is expressed as A:B
 - **A**: Segment (value **in** segment register)
 - **B**: Offset within the segment
 - Physical address = $A * 0x10 + B$
 - Segment registers:
 - CS for code, DS for data, SS for stack, ES, FS, and GS
 - Stores the base address of the segment
- Real mode is rarely used nowadays, apart from boot

Recall: x86 microprocessor modes

- Protected mode (IA-32):
 - Memory model: **Segmentation + Paging**
 - Logical address is expressed as A:B
 - A: Segment selector (Index of Global Descriptor Table)
 - B: Offset within the segment (32-bits)
 - Logical address = Segment base address (specified by GDT[A]) + B
 - The paging mechanism can be optionally enabled

Recall: IA-32 segmentation + paging



Virtual address

- **Definition:** An address that must to be translated to produce an address in physical memory
- Both logical and linear addresses are virtual

Recall: x86 microprocessor modes

- Long mode (x86-64):
 - The memory model that most of you x86-64 computers run on
 - Segmentation is almost obsolete:
 - Except for FS and GS
 - Use flat memory model (i.e., descriptor base = 0, limit = max)
 - Check more details:
 - <http://wiki.osdev.org/Segmentation>
 - <http://wiki.osdev.org/X86-64>

Managing page tables

- Linux: architecture independent (mostly)
 - Avoids segmentation (only Intel supports it)
- Abstract structures to model 4-level page tables
 - Actual page tables are stored in a machine-specific manner

Forms of fragmentation

- Fragmentation = “wasted” memory because it’s not usable due to the allocation granularity
- **External fragmentation:** the unusable memory between contiguous allocations
 - New requests for memory may not be able to find memory that is both contiguous and large enough
 - Only applicable for systems that can only allocate memory in contiguous regions (e.g., segments)
- **Internal fragmentation:** the unusable memory within the allocated blocks
 - Only applicable when allocation occurs in blocks (e.g., paging)

Summary

- Virtual memory mechanisms:
 - Paging
 - Segmentation
- Page tables, multi-level paging, PDEs/PTEs, etc.
- TLB
- Memory protection
- Memory layout with virtual memory

Demand paging

Demand paging

- Load pages into memory only as needed
 - On the first access
 - Pages that are never used never get loaded
- Use **valid bit** in the page table entry:
 - Valid implies
 - The page is in memory (i.e., valid mapping)
 - Invalid implies either:
 - Valid mapping, but the page is not in memory
 - Illegal memory access (e.g., null pointer dereference)
- OS should do something more for the invalid case to support on demand paging

Demand paging: At process start

- Open executable file, and parse executable headers
- Set up memory mapping (stack, text, data, bss, etc.)
 - Set up internal virtual memory mapping for OS
 - But don't load anything yet
- Load first page (entry point) and allocate initial stack page
 - Set up the page table (for CPU/MMU)
- Run the program

Memory mapping

- Executable files and libraries must be brought into process's virtual address space
 - File is mapped into the process's memory
 - As pages are referenced, page frames will be allocated and pages are loaded into them

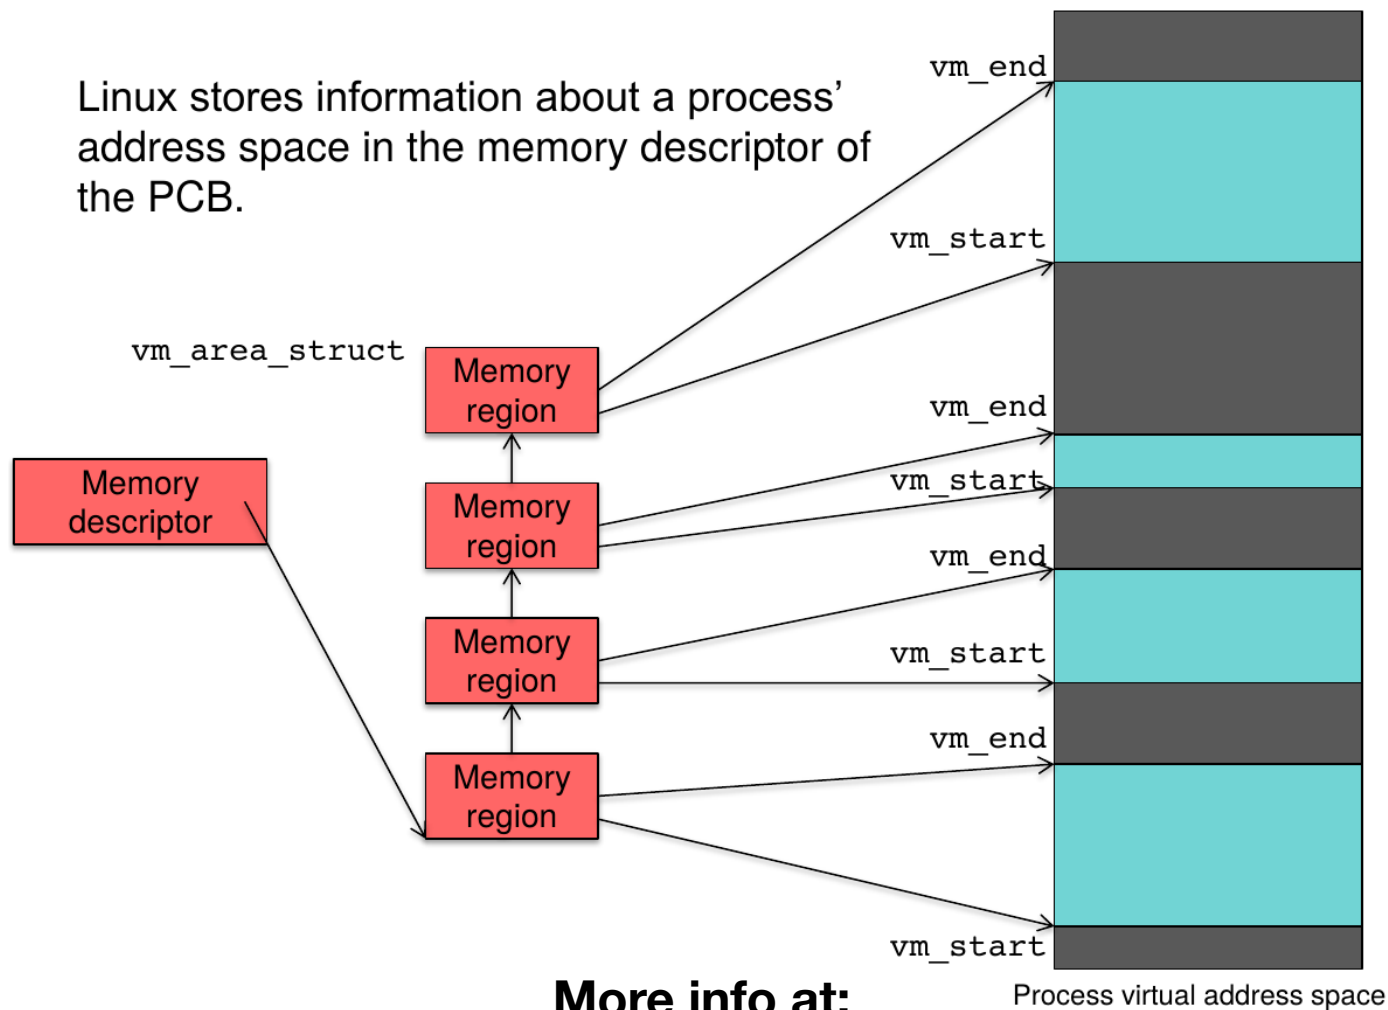
Memory mapping

- ***vm_area_struct*** in Linux:
 - Define true semantics of virtual memory space
 - This is not a page table
 - Page table only shows the partial semantics of virtual memory space
 - Start of VM region, end of region, access rights
- Several of these are created for each mapped image
 - Executable code, initialized data, uninitialized data

Demand paging: page fault handling

- Eventually the process will access an address without a valid page:
 - OS gets a page fault from the MMU
- What happens?
 - Kernel searches a tree structure of memory allocations for the process to see if the faulting address is mapped in VA space
 - If not mapped, send a SEGV signal to the process
 - If mapped:
 - Check if the type of access is valid for that page
 - If valid, the page is not in the memory (but in the file) so the OS fetches it from the file

Keeping track of a processes' memory region



More info at:
https://linux-kernel-labs.github.io/master/labs/memory_mapping.html

Demand paging components

- Basic mechanism to intercept page accesses
- Swap medium
- Page replacement mechanisms and policies

Page replacement

- OS may need to kick out a page from memory
 - Especially if running out of physical memory space
- When replacing a page (i.e., when evicting a page from memory)
 - If the page came from a file and was not modified
 - Discard, we can always get it back from the file
 - If the page is dirty, it must be saved in a page file (aka swap file)
 - Page file: a file (or disk partition) that holds data of memory pages
 - If the page is not backed up by a file
 - Always save into a page file

Swap medium

- Swap medium in practice
 - Windows: `pagefile.sys`
 - Linux: swap partition or swap file
 - OS X: multiple swap files in “`/private/var/vm/swapfile`”

Demand paging: getting a page

- The page we need is either in the mapped file (executable or library) or in a swap file
- Read page into physical memory:
 - Find a free page frame (evict one if necessary)
 - Read the page from the file (slow!)
 - Update the page table for the process
 - Resume the process at the instruction that faulted

Cost

- Handle page fault exception: 400 usec
- Disk seek and read: 10 msec
- Memory access: 100 ns
- Page fault degrades performance by around 100,000x
 - Avoid page faults if (or as much as) possible!!
 - We need a good replacement policy for good performance
- Q: Performance impacts of a program exhibiting
 - Sequential memory accesses
 - Random memory accesses

FIFO replacement

- First in, first out
- Good: May get rid of initialization code or other code that's no longer used
- Bad: May get rid of a frequently accessed page

Least Recently Used (LRU)

- Timestamp a page when it is accessed
- When we need to remove a page, search for the one with the oldest timestamp
- Good properties but:
 - Timestamping is costly...
 - We can't do it with most MMUs

Not Frequently Used Replacement

- Approximates LRU behavior
 - Each PTE has a reference bit
 - CPU sets the reference bit on memory reference
 - If it's not set then the frame hasn't been used for a while
 - Keep a counter for each page frame
 - At each clock interrupt:
 - Add the reference bit of each frame to its counter
 - Clear reference bit
 - To evict a page, choose the frame with the lowest counter
- Problem:
 - No sense of time: a page that was used a lot a long time may still have a high count
 - Updating counters is expensive

Clock

- Arrange physical pages in a logical circle (circular queue)
 - Clock hand points to first frame
- Paging hardware keeps one reference bit per frame
- On page fault:
 - Check reference bit of the page frame where clock hand points to
 - If 1, it's been used recently:
 - Clear the reference bit and advance
 - If 0, evict this page

Enhanced clock using modify bits

- Use the reference and modify bits of the page together
- Choices for replacement (reference, modify):
 - (0, 0): not referenced or modified recently
 - Good candidate for replacement
 - (0, 1): not referenced but modified
 - The page will have to be saved before replacement
 - (1, 0): recently used
 - Less ideal: a program is likely to use it again
 - (1, 1): recently used and modified:
 - Least ideal: a program may use again AND we'll have to save the page to swap if we replace it

CODE

- How to measure (empirically using software) cache properties?
- How to cause TLB misses deliberately?

Summary

- On demand paging:
 - Lazy loading of executables
 - Simulating/virtualizing more memory than the available physical RAM
- Intercepting page accesses using the valid bit
- Trapping on the page faults