

# More virtualization and TEEs

CS503: Operating systems, Spring 2019

Pedro Fonseca  
Department of Computer Science  
Purdue University

# Admin

- Ask for office hours if necessary (e.g., send email)
- Online students: also make sure you talk with TAs and me

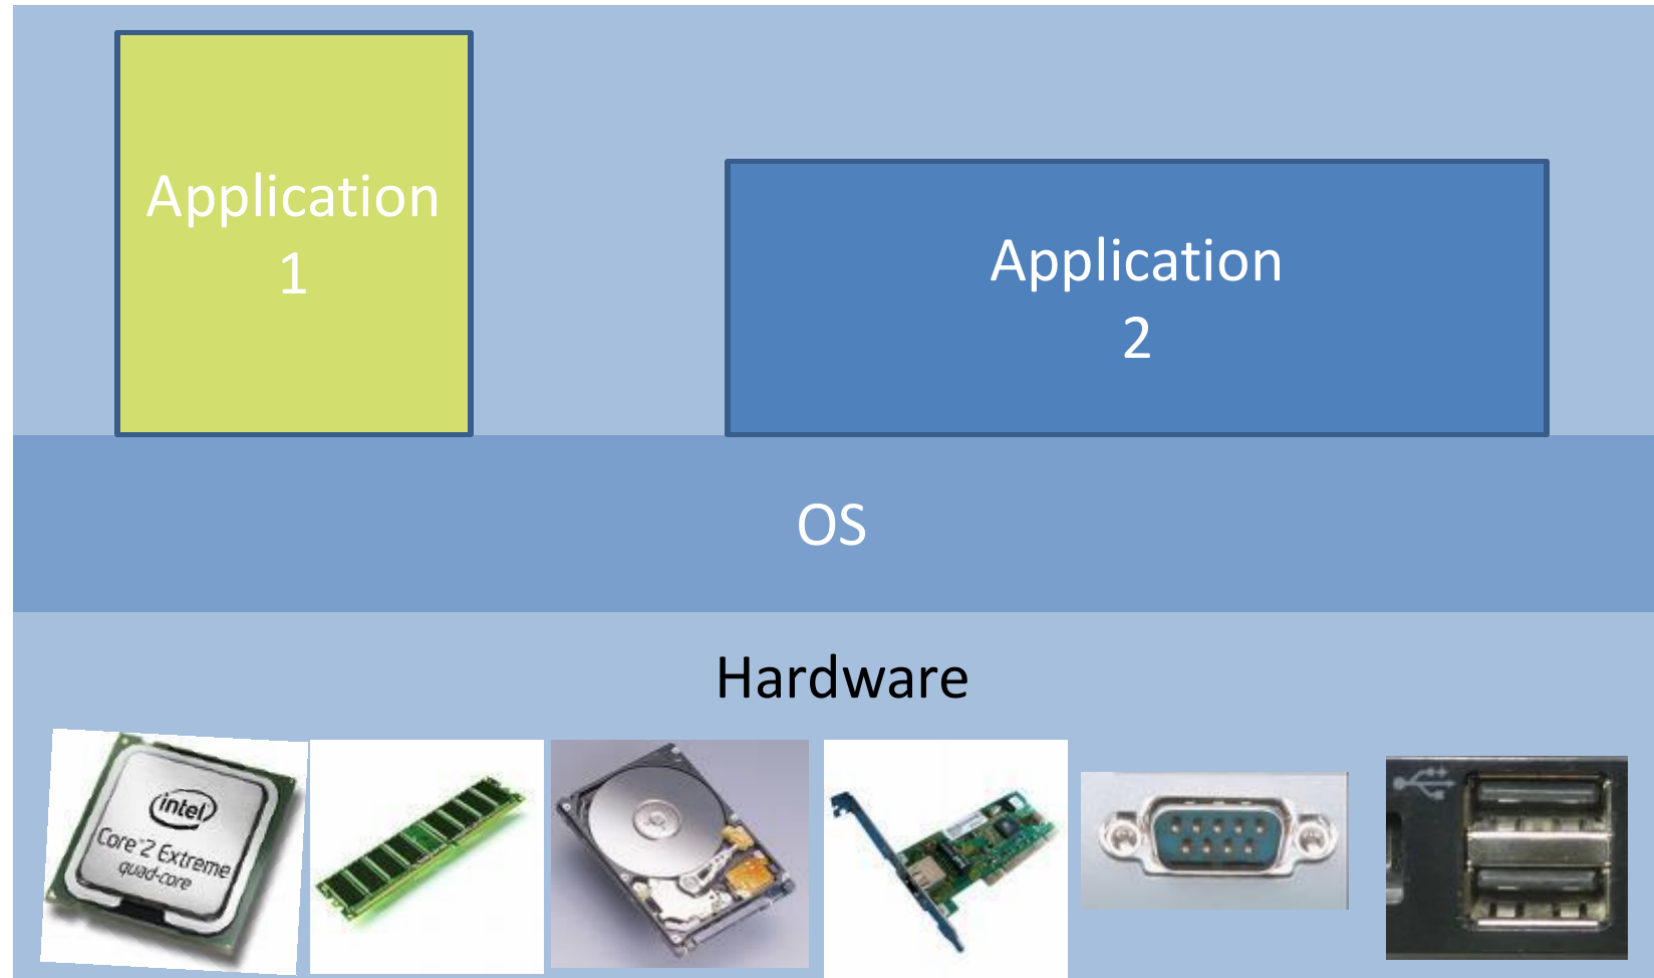
# Next few lectures

- Advanced topics:
  - Virtualization / containers / SGX
  - Linux
  - Systems research
  - Operating systems + distributed systems
  - Review for exam


# Last lecture

- VM-level virtualization
- Challenge in multiplexing
- Para-virtualization
- Advanced OS features
  - Checkpoint and restart
  - Migration

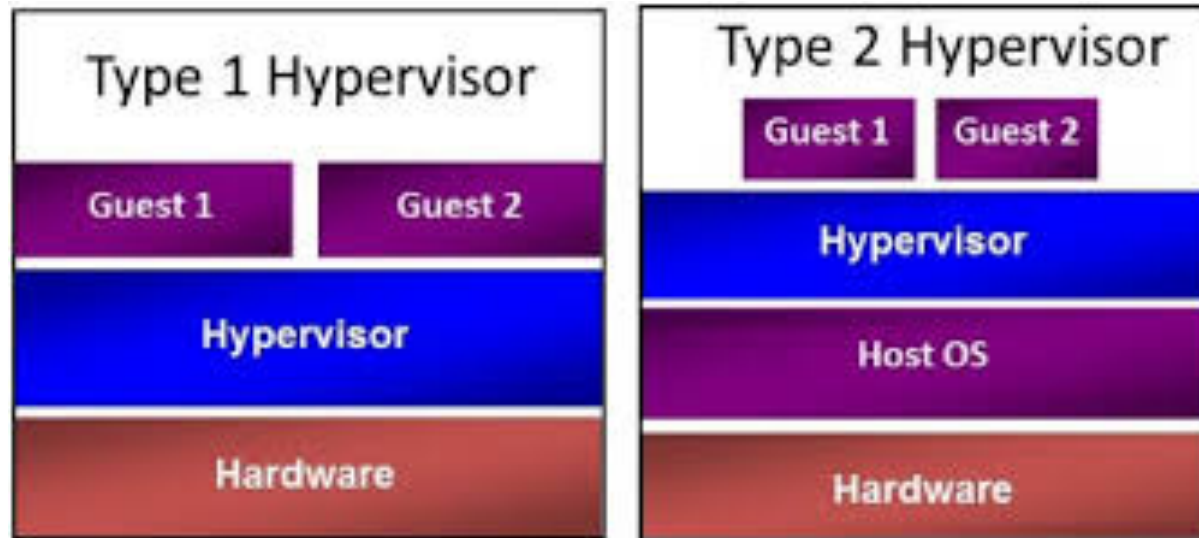
# Recall: Traditional



# Recall: VM virtualization



# Type 1 vs. type 2 hypervisors



e.g., Xen

e.g., VirtualBox,  
QEMU/KVM

**What are the advantages of each?**

# Recall: Advanced topics and discussion

- Scheduling multiple VMs
  - How does this affect the VMM?
  - Impact of caches
- Suspend and resume
- Migration of VMs
  - What are the uses?
  - What are the challenges?

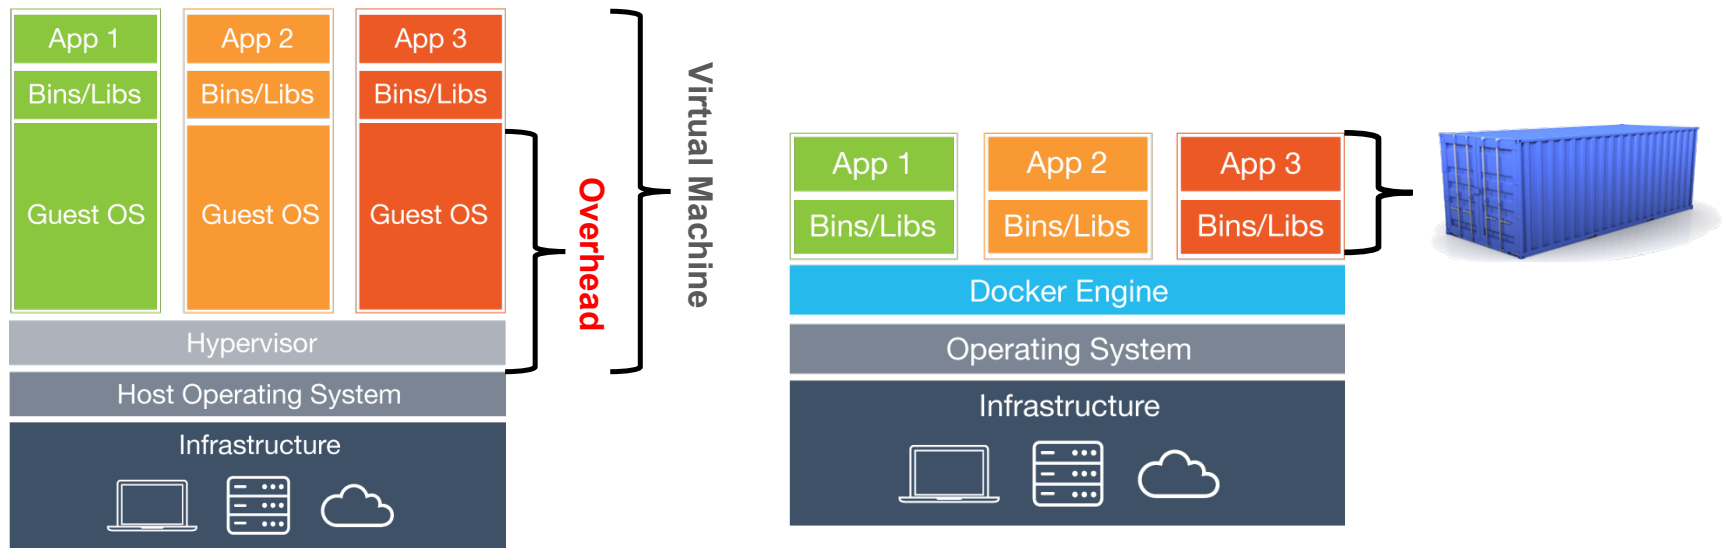


# Process-level virtualization (containers)

# Process-level virtualization

- Also known as *OS-level virtualization*:
  - Used to implement the *container* abstraction
- Run application in an isolated environment

# VMs vs. containers



- Containers are generally “lighter” than VMs

|   | Container                                      | Virtual Machine                     | Bare-Metal x86 Server          |
|---|--|-------------------------------------|--------------------------------|
| Underlying Platform   | OS on Virtual Machine or Bare-Metal x86 Server | Hypervisor on Bare-Metal x86 Server | N/A                            |
| Performance: Speed and Consistency  | Average  | Average                             | Fastest                        |
| Provisioning Time   | Seconds  | Minutes                             | Hours                          |
| Tenant Isolation Enforcement  | OS Kernel                                      | Hypervisor                          | Physical                       |
| Ideal Application Types   | Mode 2   | Mode 1 or Mode 2                    | Mode 1 or Mode 2               |
| Configuration and Reconfiguration Flexibility   | Highest  | Medium                              | Lowest                         |
| Host Consolidation Density  | Maximum  | Average                             | None                           |
| Application Portability   | Application Packaging/Manifest*                | VM Image, VM Migration Tools        | Backup and Restore, ISO Images |
| Granularity   | Extremely Small                                | Average                             | Largest                        |
| <p>*While application portability is somewhat easier in container environments that are leveraging a container management and orchestration solution, portability should not assumed to be universal — differences in the underlying host OS below the containers could still present some interoperability challenges.</p> |  |                                     |                                |

Source: Gartner (September 2015)

**\* Take this table with a grain of salt**

# Containers

- Recent trend to develop and deploy applications:
  - Ensures the same environment
  - Sandbox application / project
  - Automatically deals with dependencies
- Containers usually rely on
  - OS services to implement isolation
  - A library that implement a common environment
  - Tools to manage containers

# Using containers with Docker

- Write a “dockerfile”:

```
FROM ubuntu:18.04
COPY . /app
RUN make /app
CMD python /app/app.py
```

- Build an “image” (layered read-only FS)
- Deploy the image on a machine
- Other container management systems: rkt, LXC

Q: And how to use a VM?

[https://docs.docker.com/develop/develop-images/dockerfile\\_best-practices/](https://docs.docker.com/develop/develop-images/dockerfile_best-practices/)

# Discussion

- What are the (minimum) OS services required to implement containers?
- Container orchestration (e.g., Kubernetes)
- When are VMs better?
- Are there security differences
- How to address the network isolation? Would it be useful?

# Trusted execution environments



# Security issues with the cloud

- Future of clouds: all your data will be in clouds
  - Email
  - Location-based service
  - Bitcoin
  - Password management service
  - DNA matching service
- Q. Should you trust clouds?
  - Should you trust Amazon when using Amazon EC2?
  - Should you trust Google when using gmail?

# Trusted execution environments

- Intel SGX, ARM TrustZones
- Implement the secure **enclave** abstraction
- Provide several properties to the code running inside the enclave

# Properties: confidentiality

- An adversary outside of the enclave cannot inspect the state of execution inside the enclave
- Even if the OS is compromised or execute a malicious application on the machine

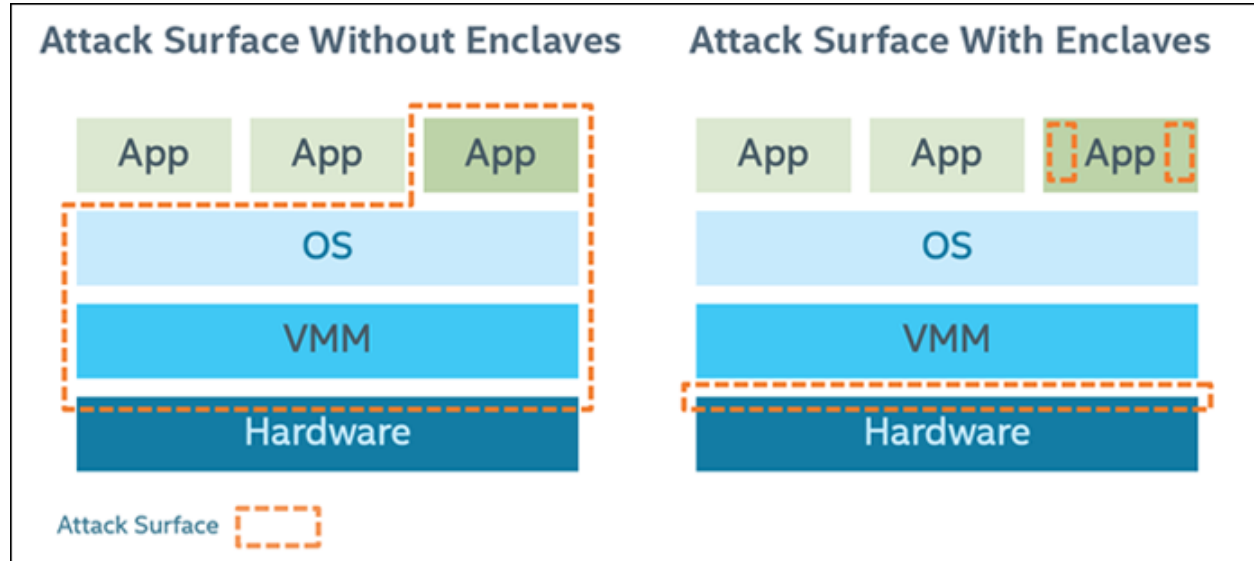
# Properties: integrity

- Guarantees that an attacker outside of the enclave cannot affect the result of computation inside the enclave:
  - Except from supplying inputs through the defined interface
- Ensures correctness of the computation running inside the enclave even if the OS

# Properties: attestation

- Provides an (unforgeable) proof that enables a remote party to verify what has run inside the enclave
- Even if they don't have physical access to the machine
- Relies on cryptography (and requires trusting the HW)

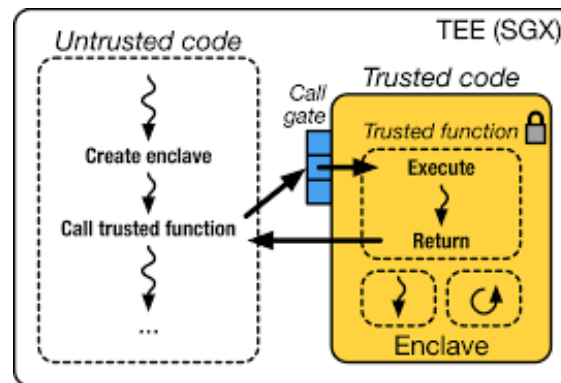
# SGX attack surface



- Requires special hardware
- Hardware-support can also prevent physical attacks by reading the memory bus

Q: SGX + hypervisor makes sense?

# Secure computation with SGX enclaves



# Discussion topics

- Hows does SGX compare with user/kernel protection?
- SGX + hypervisors makes sense
- What cryptography is necessary? Where are keys stored?
- Any important properties not guaranteed by SGX?
- Practical security
- Mathematical alternatives to HW-based solutions?



# Summary

- Hypervisors
  - Type 1 vs. type 2
  - Advantages
- OS-level / process-level virtualization
  - Container abstraction
  - Light virtualization
  - Performance, isolation, security
- Trusted execution environments